

NEF and Nengo

Centre for Theoretical Neuroscience
Nengo Summer School

Chris & Terry



Three views

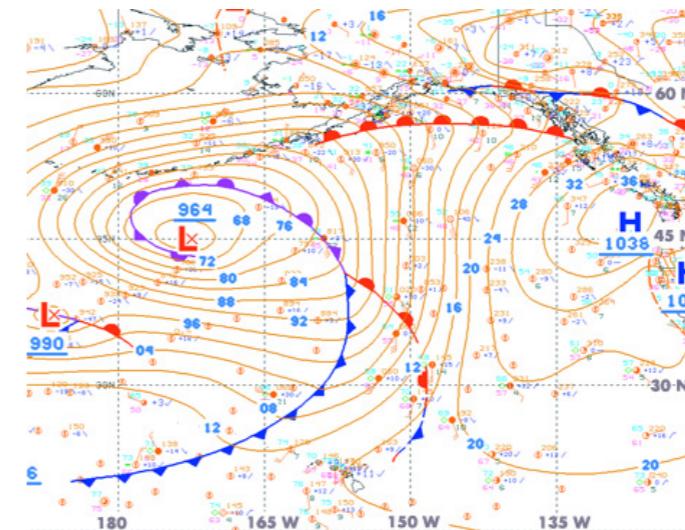
- Symbolicism



- Connectionism

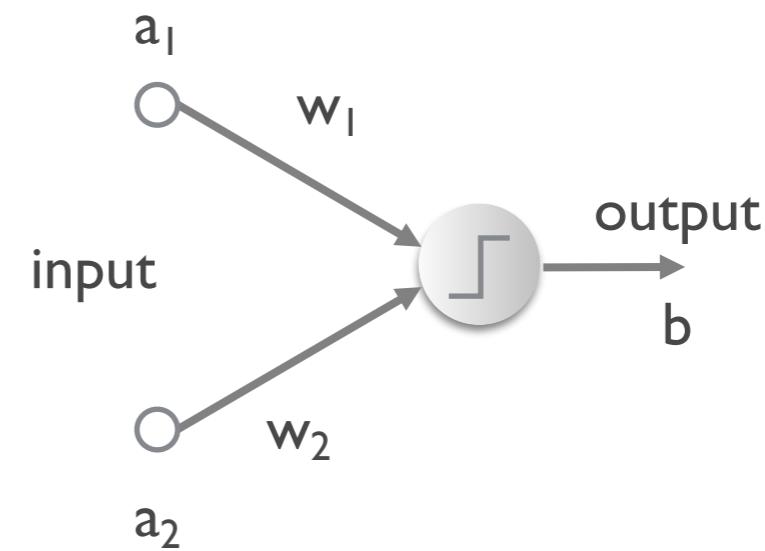
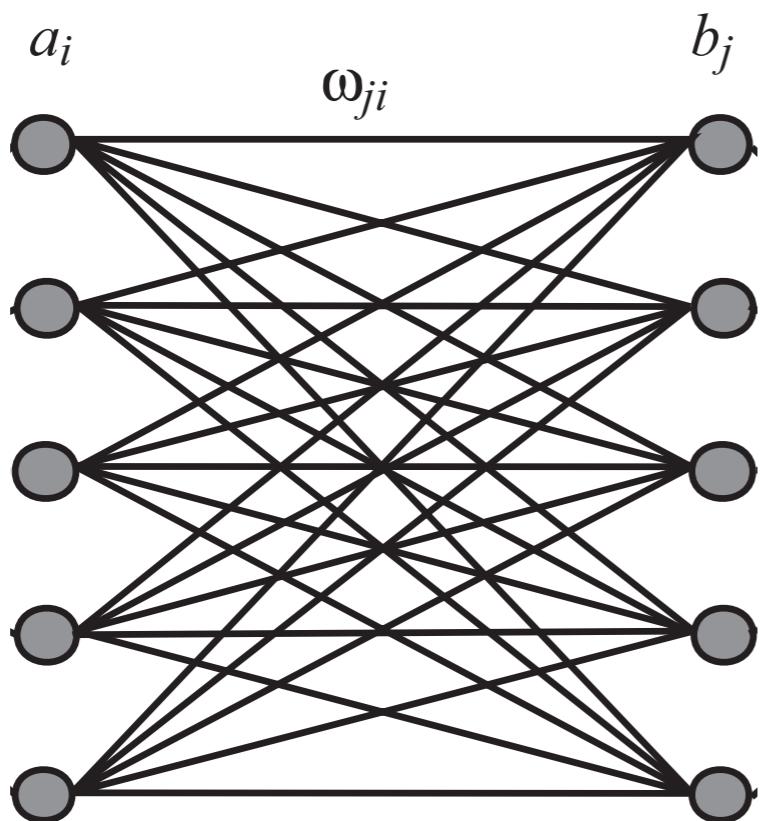


- Dynamicism



Connectionism

- ‘Brain-like’ networks



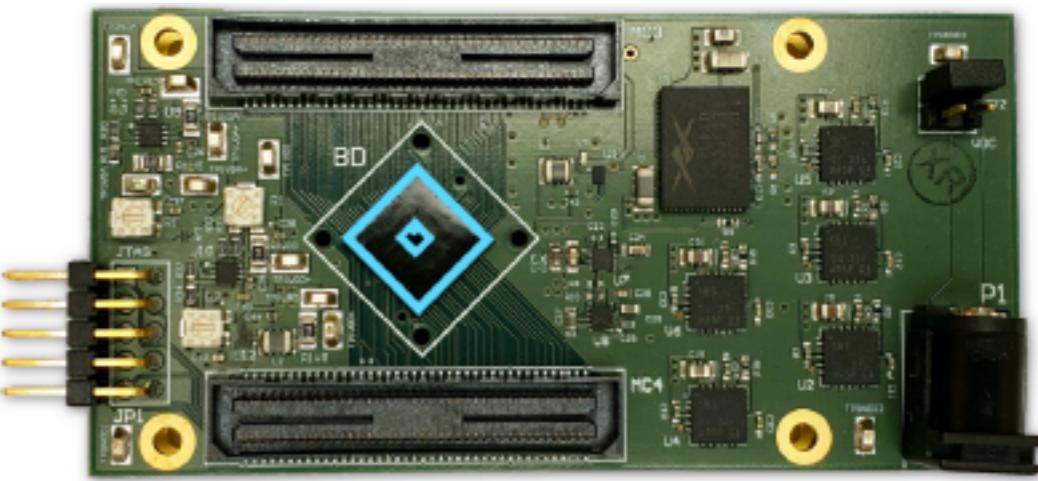
Deep Learning

- Backpropagation is an effective way to pick weights
- Challenges:
 - Relating to symbolism (e.g., explainability, representation of concepts, high-level reasoning)
 - Incorporating time (e.g., long time scales, online learning, latency)

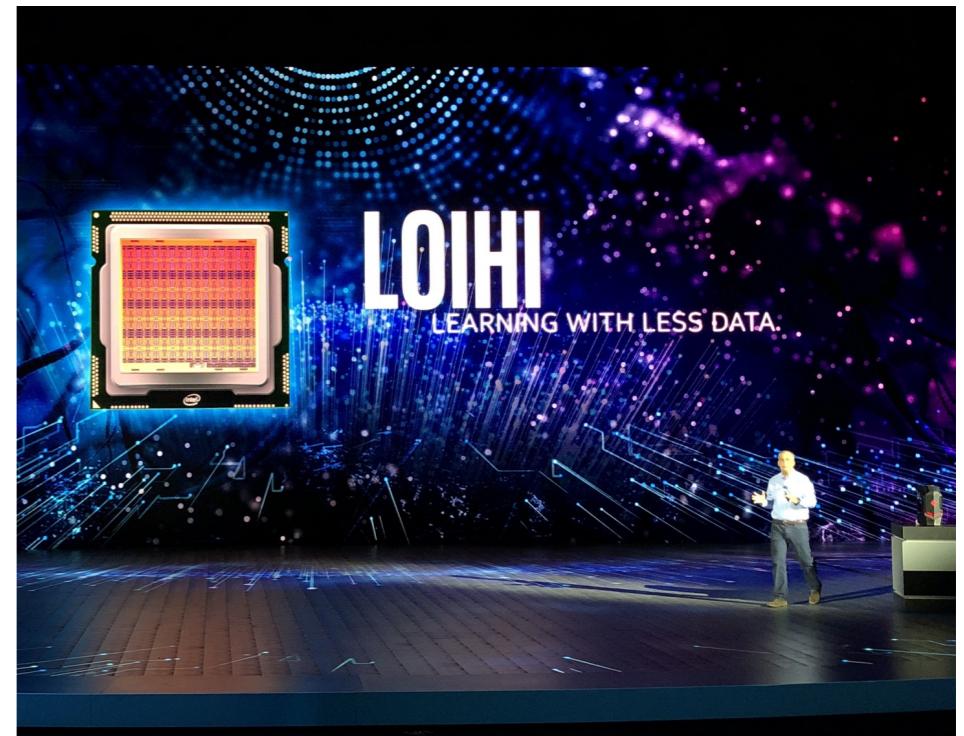
Neuromorphics

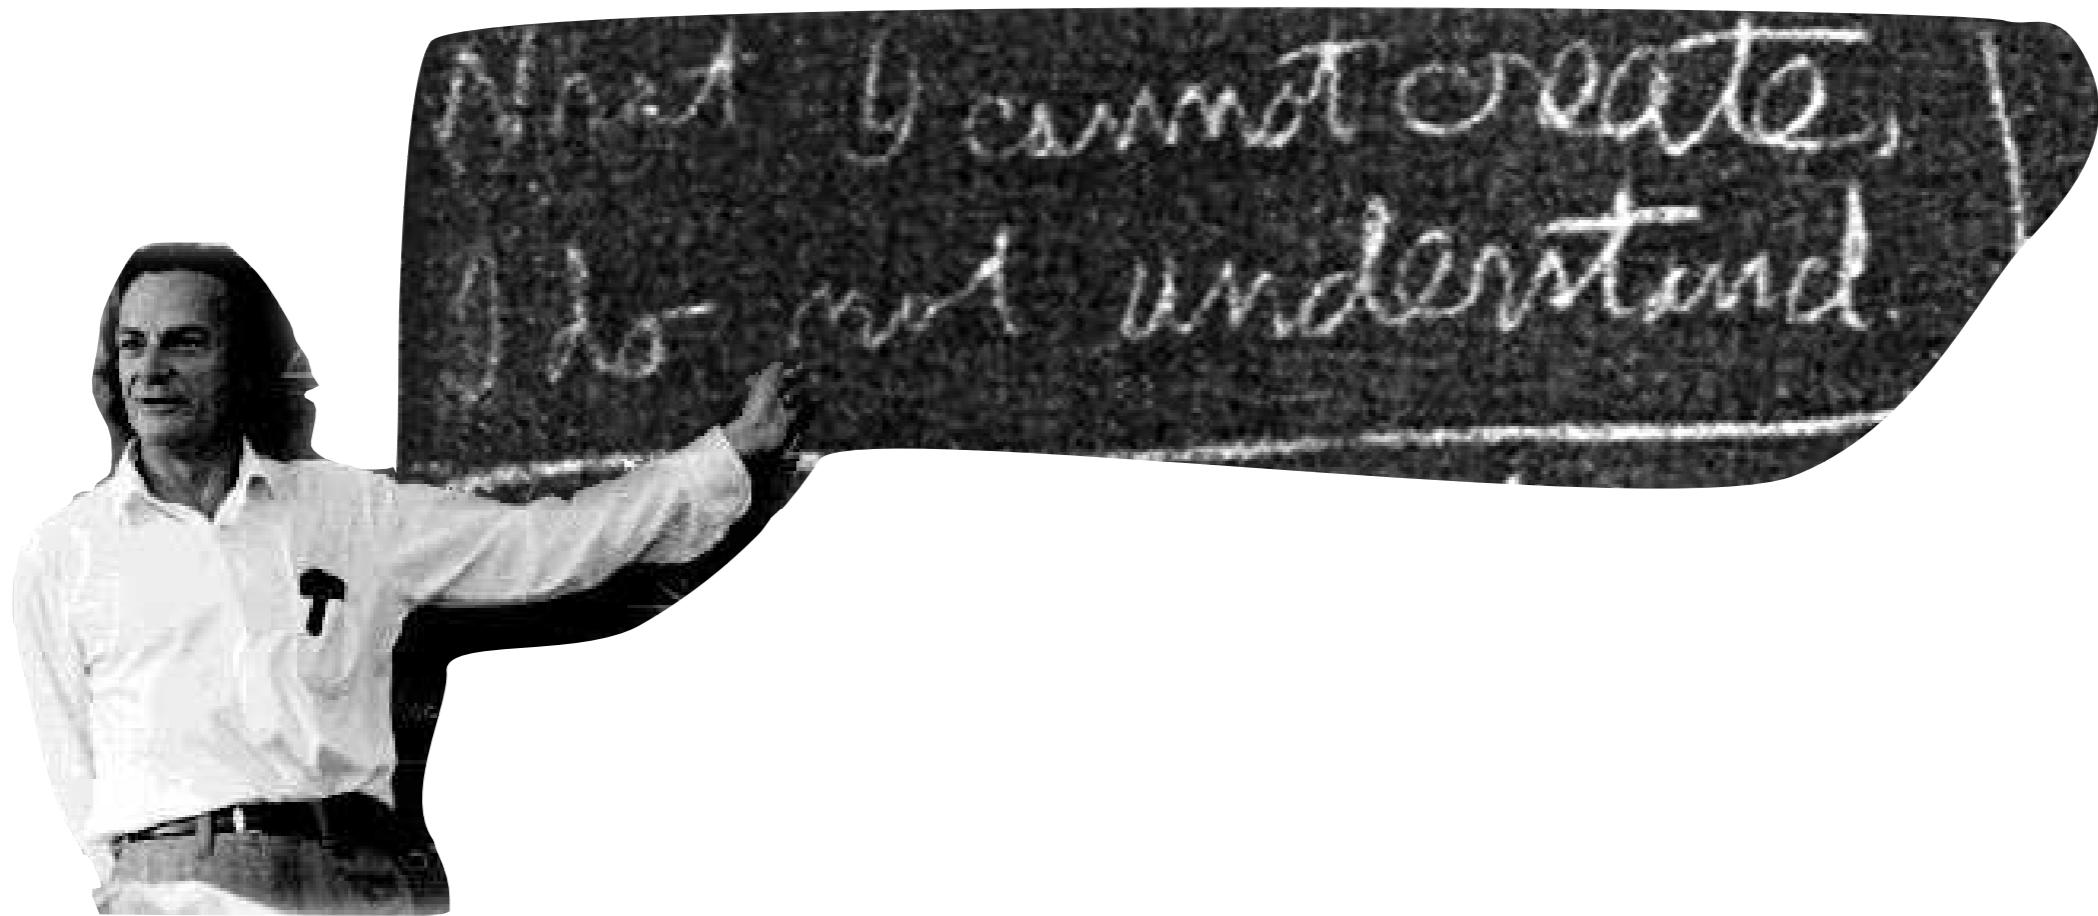
- Improve neural computing:
 - Lower power, real-time, online learning, robust
- Advance cognitive computing:
 - Large-scale neural modelling
 - Practical real-world, dynamic, and scalable AI

Braindrop



Loihi





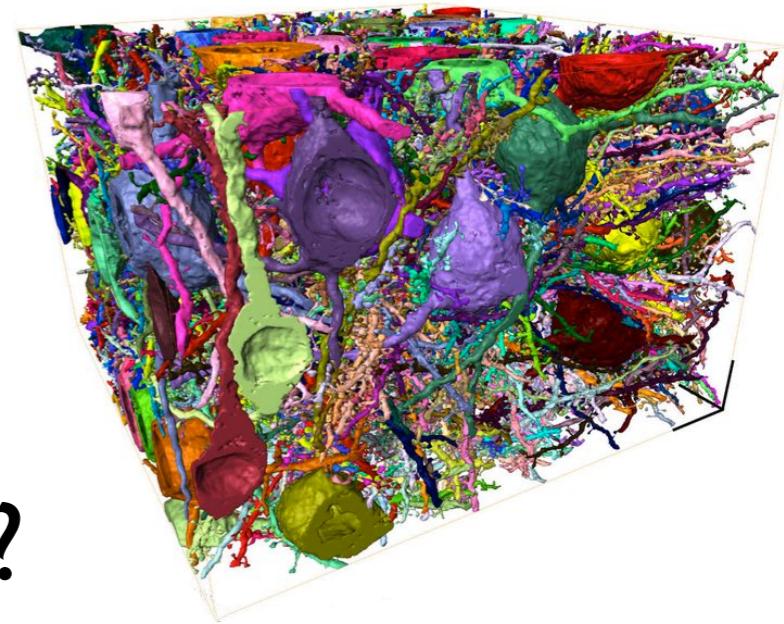
large-scale
complete behaviour

A more synthetic approach to neuroscience (to complement the analytic approach)

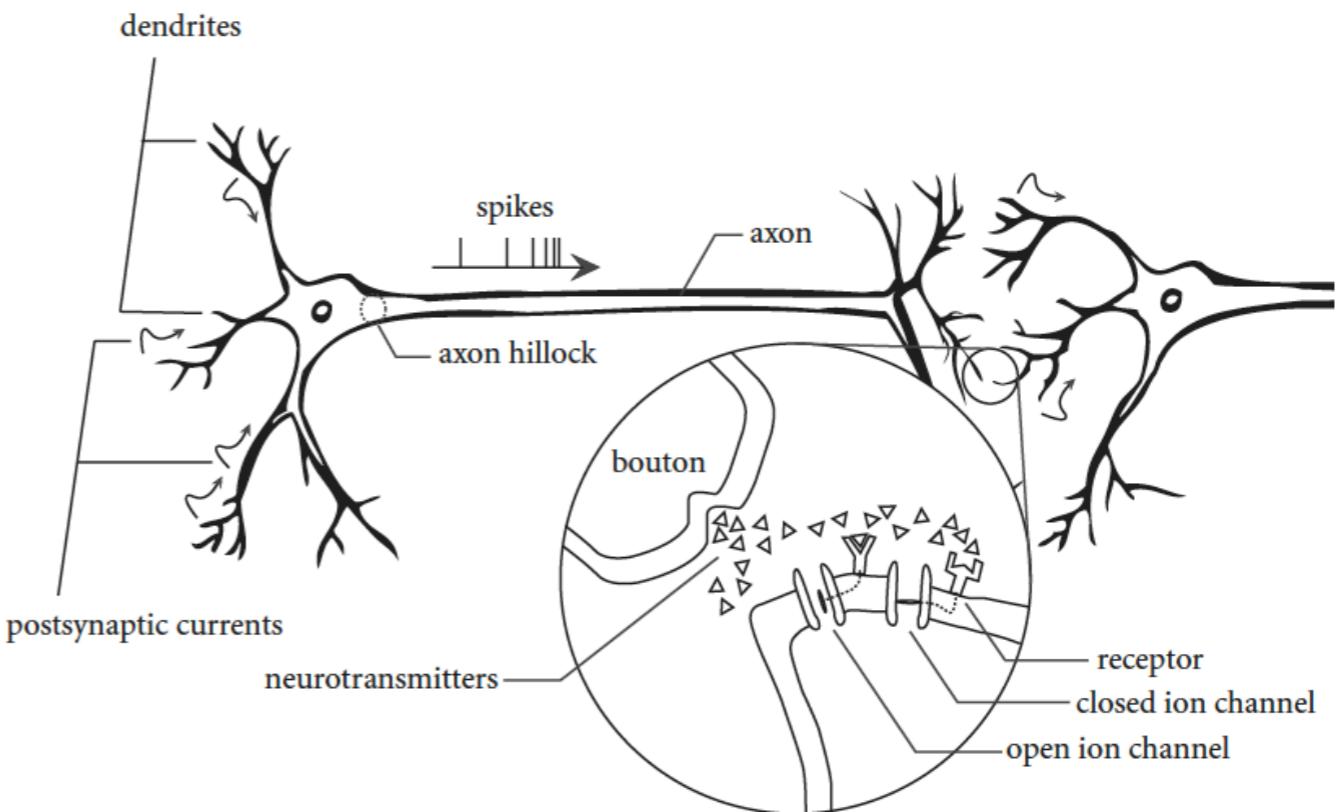
Challenges

- Include:
 - heterogeneity,
 - noise/reliability,
 - nonlinearities,
 - scaling

Neurons



- What are they good at computing?
 - Low precision, time-based, adaptive
 - Specific kinds of nonlinear functions
- How much complexity do we need?
 - As much as is useful for the questions we are asking



The NEF

Neural Engineering Framework (NEF)

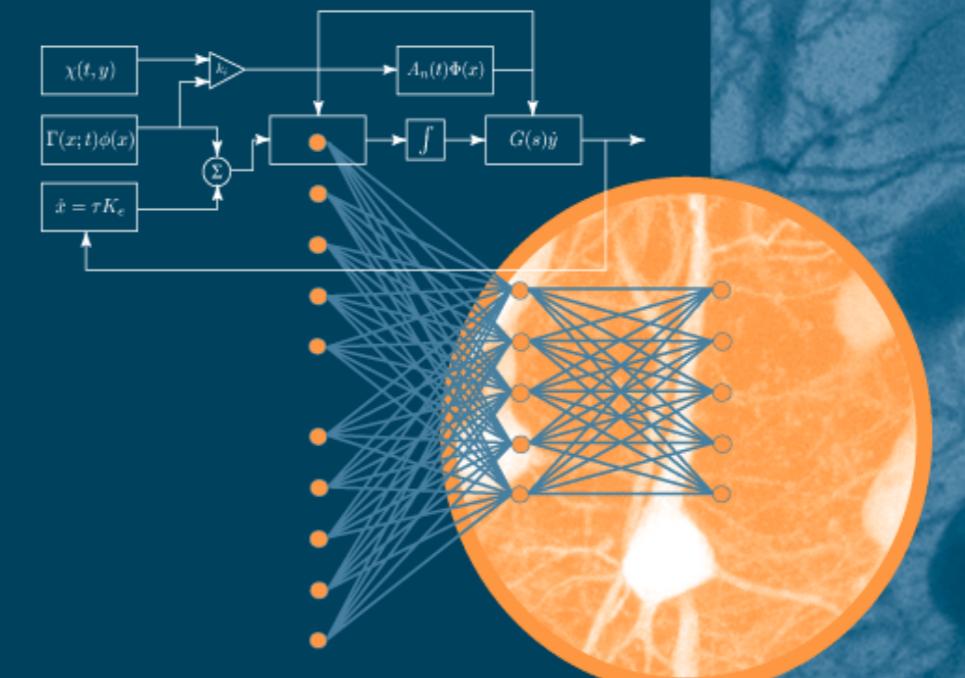
- From high-level fcn to multi-layer spiking networks
- Three principles
 - Representation
 - Computation
 - Dynamics

Eliasmith & Anderson (2003) MIT Press

Neural Engineering

COMPUTATION, REPRESENTATION, AND DYNAMICS
IN NEUROBIOLOGICAL SYSTEMS

Chris Eliasmith and Charles H. Anderson



Three Principles of the NEF

1. Representation

- Neural spike trains are nonlinear *encodings* of vector spaces that can be linearly *decoded*.

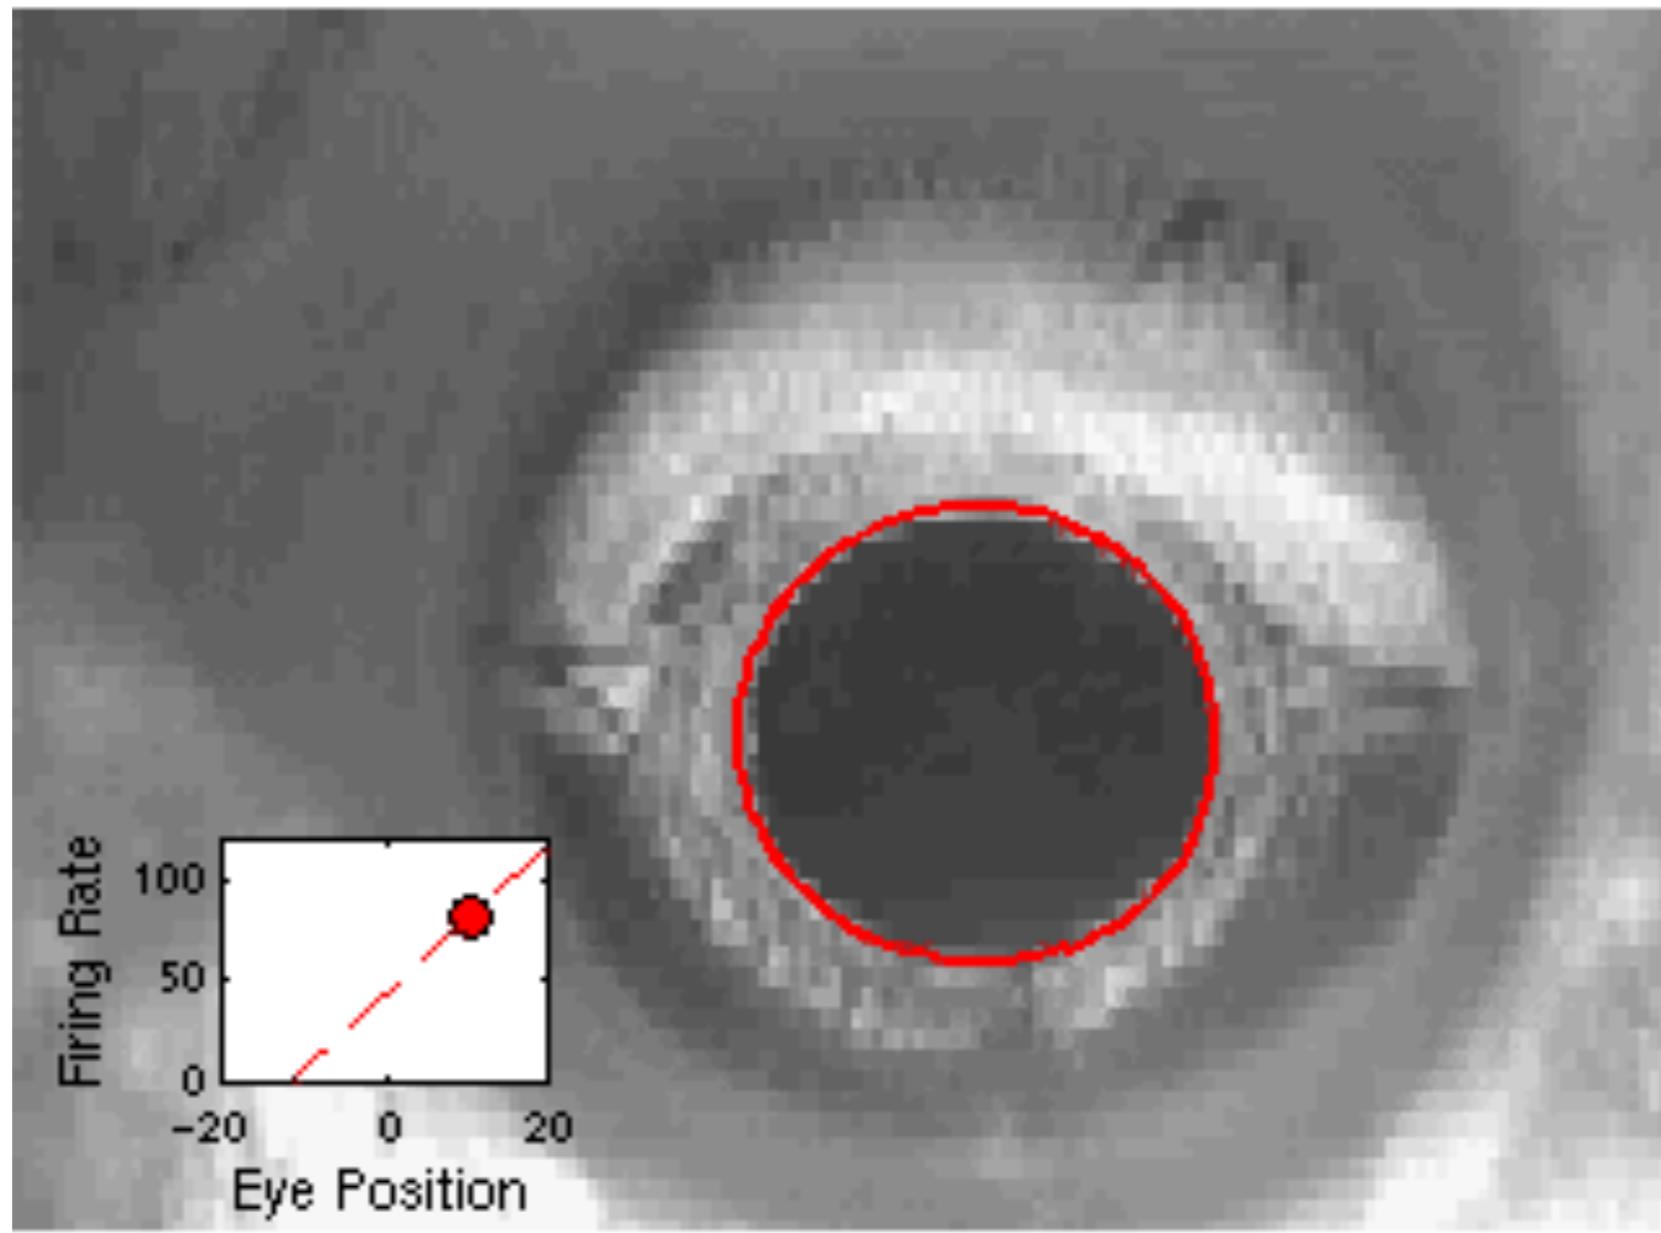
2. Computation/Transformation

- Alternate linear decodings of those encodings can compute arbitrary vector functions.

3. Dynamics

- Neural representations (1) are control theoretic state variables in a nonlinear dynamical system (2)

Principle I: Representation



Principle I: Representation

Bob

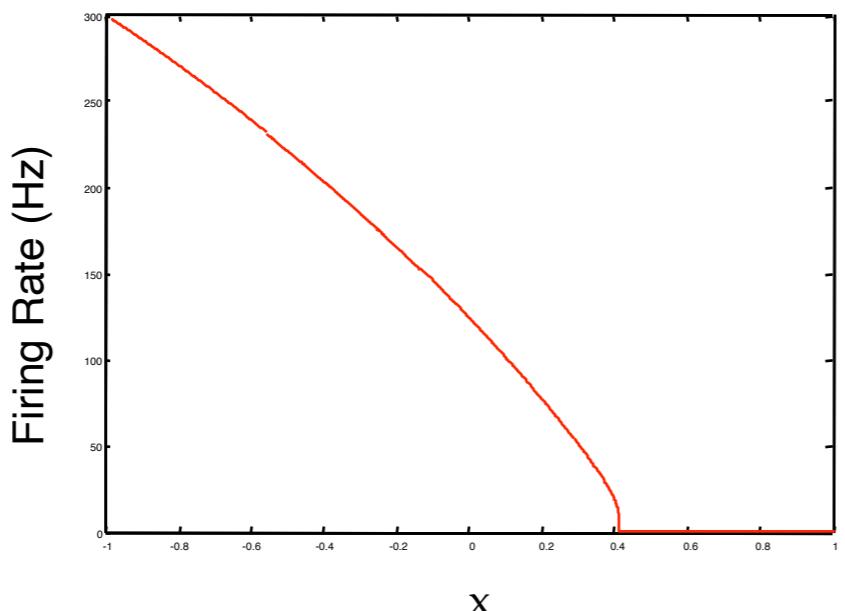
axon

Alice

axon

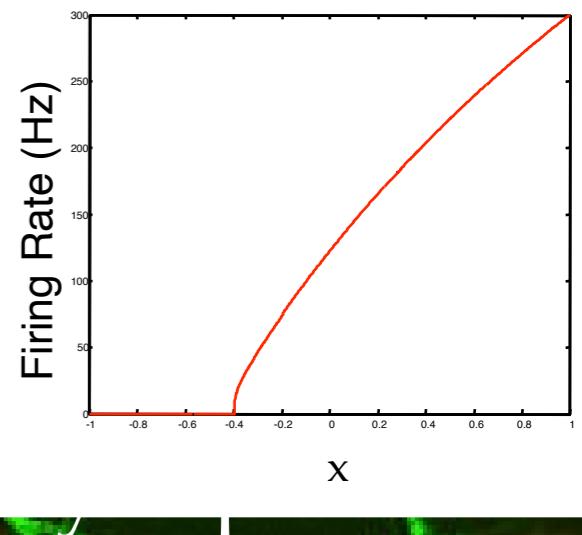
B

Neuron Tuning Curve



A

Neuron Tuning Curve



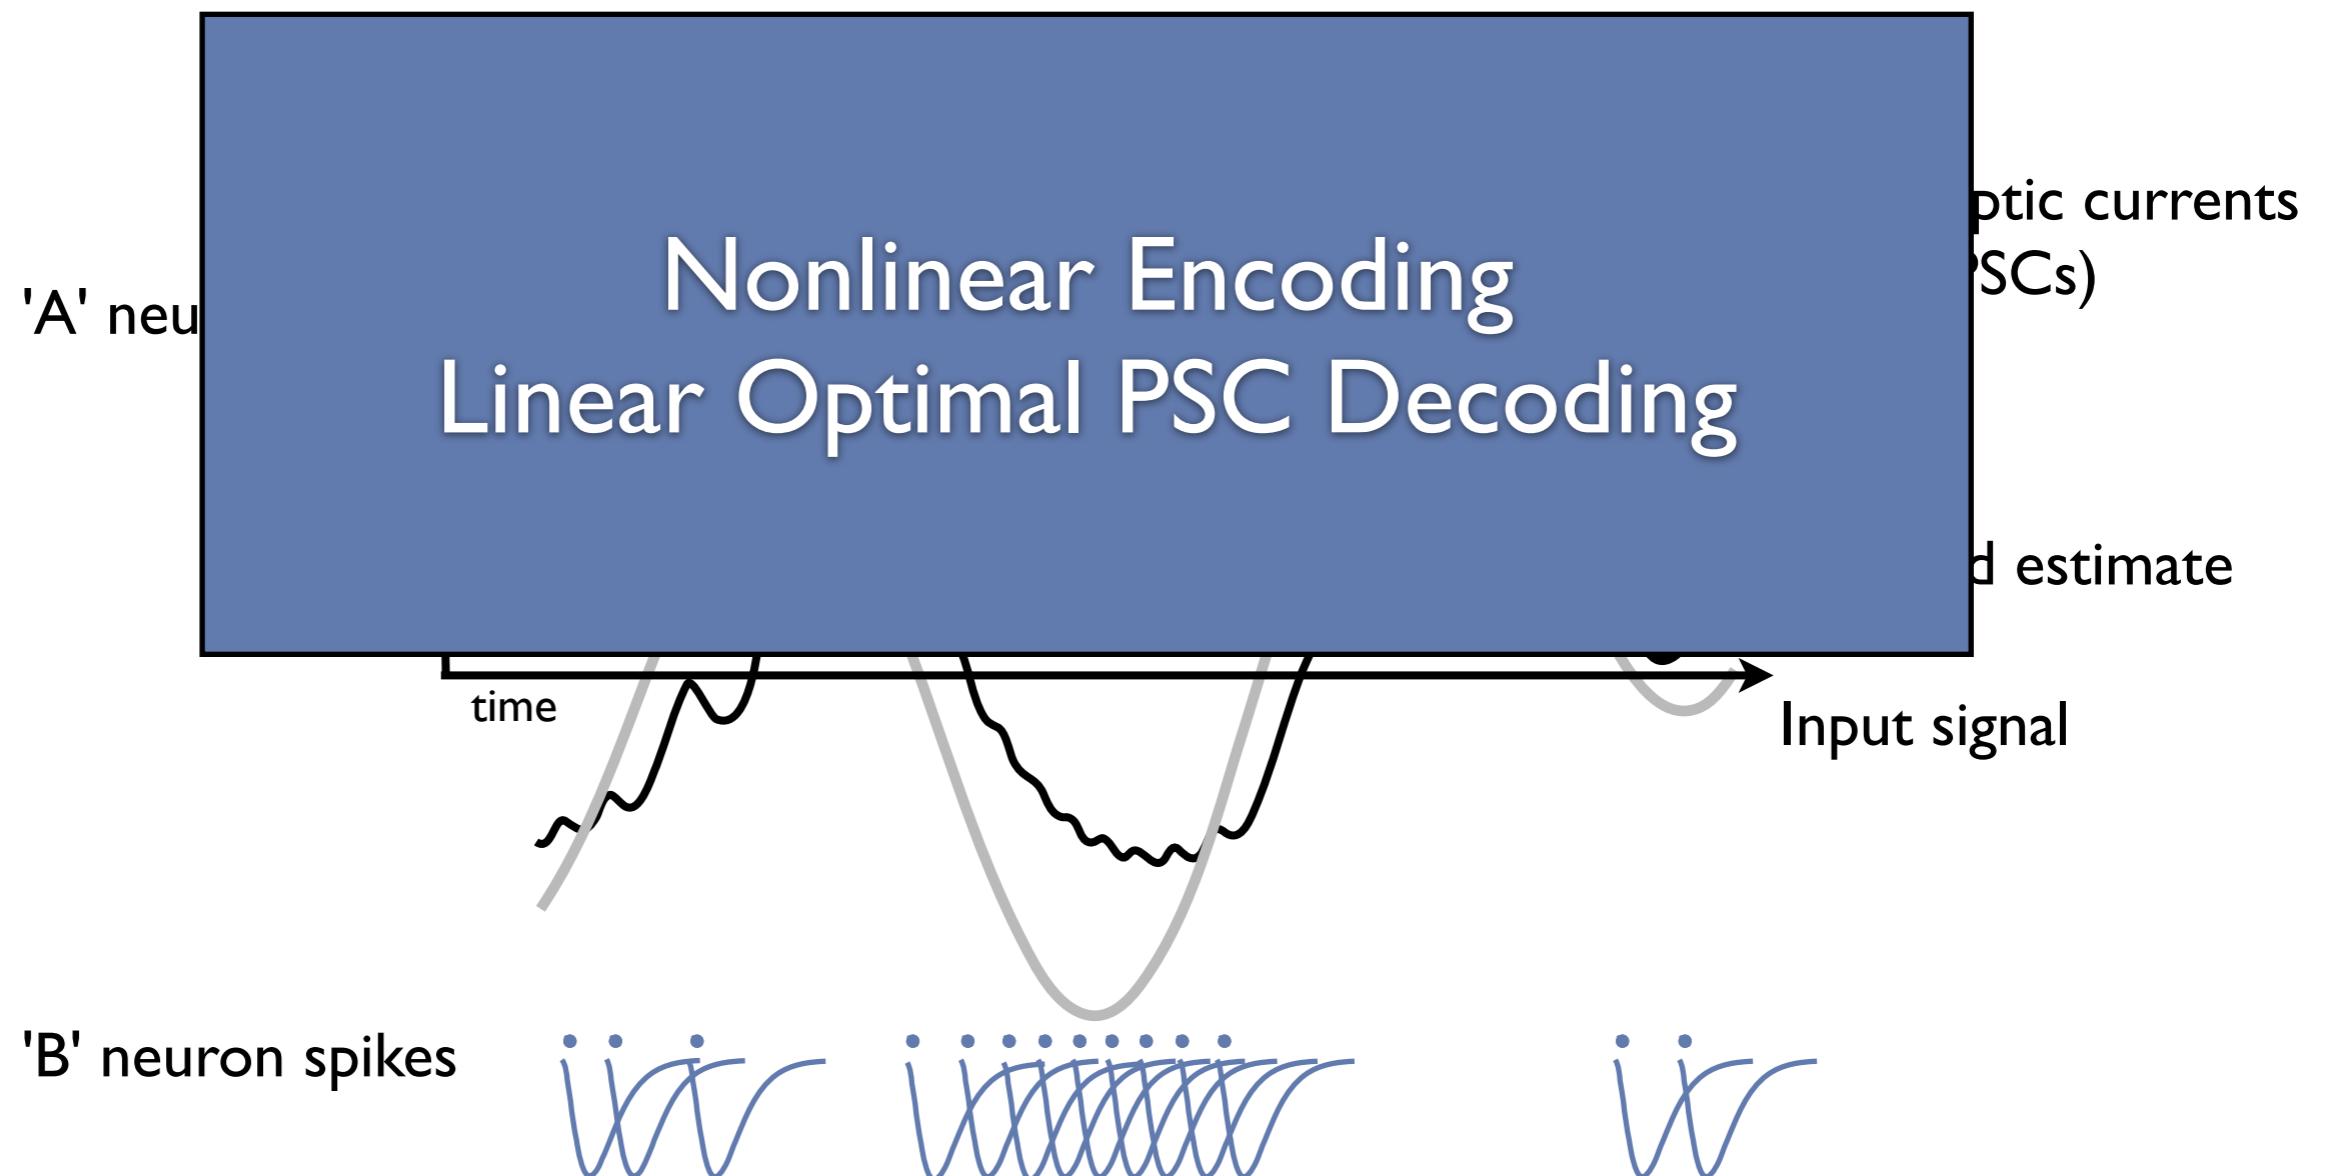
Charlie

axon

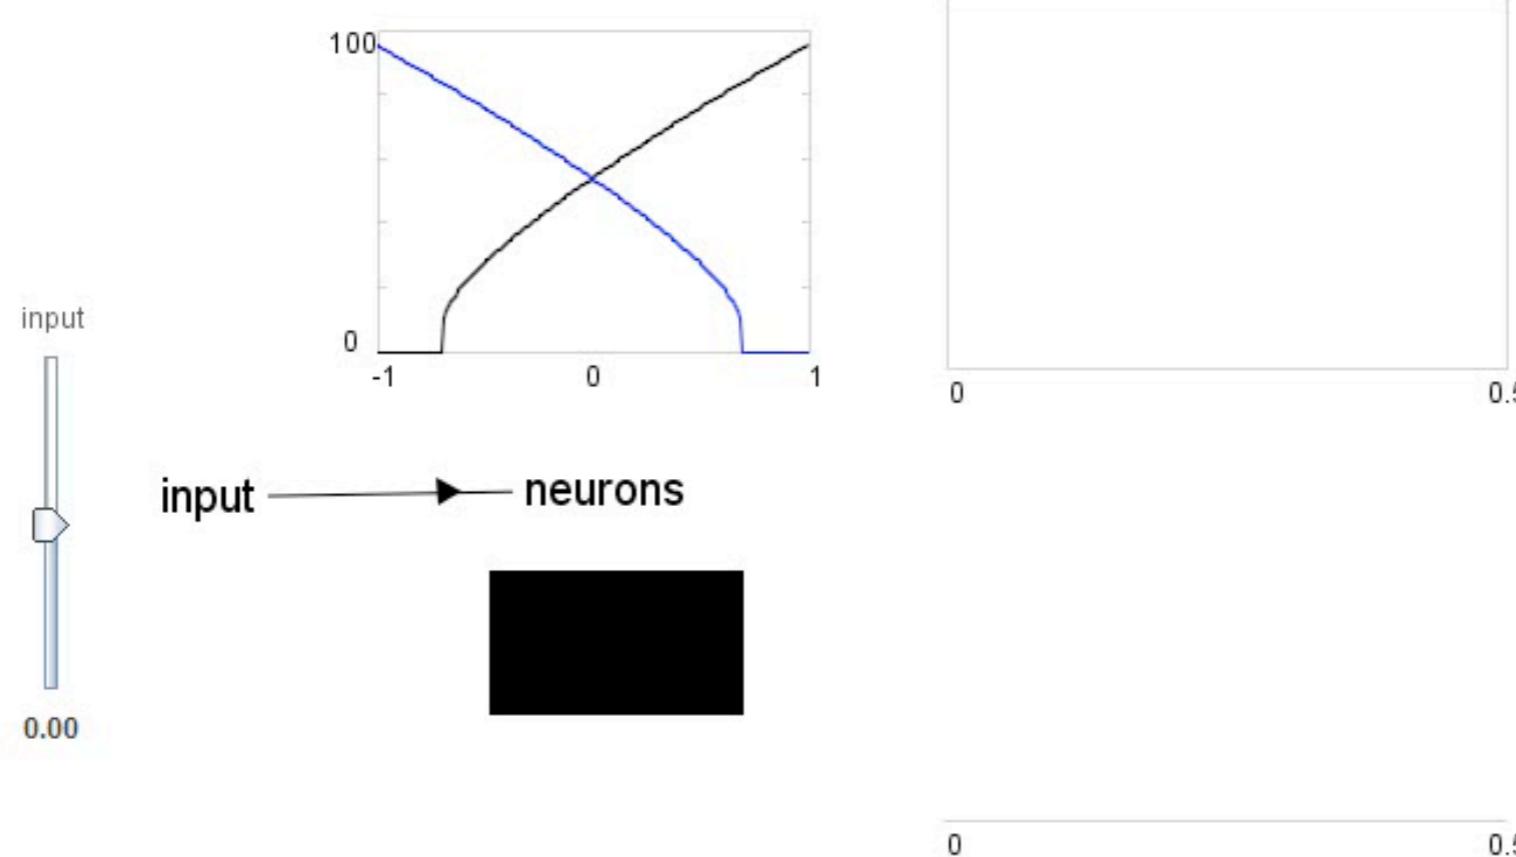
What could Charlie know about x ?

Principle I: Representation

- Encoding (stimulus \rightarrow spikes; A & B)
- Decoding (spikes \rightarrow stimulus; C)



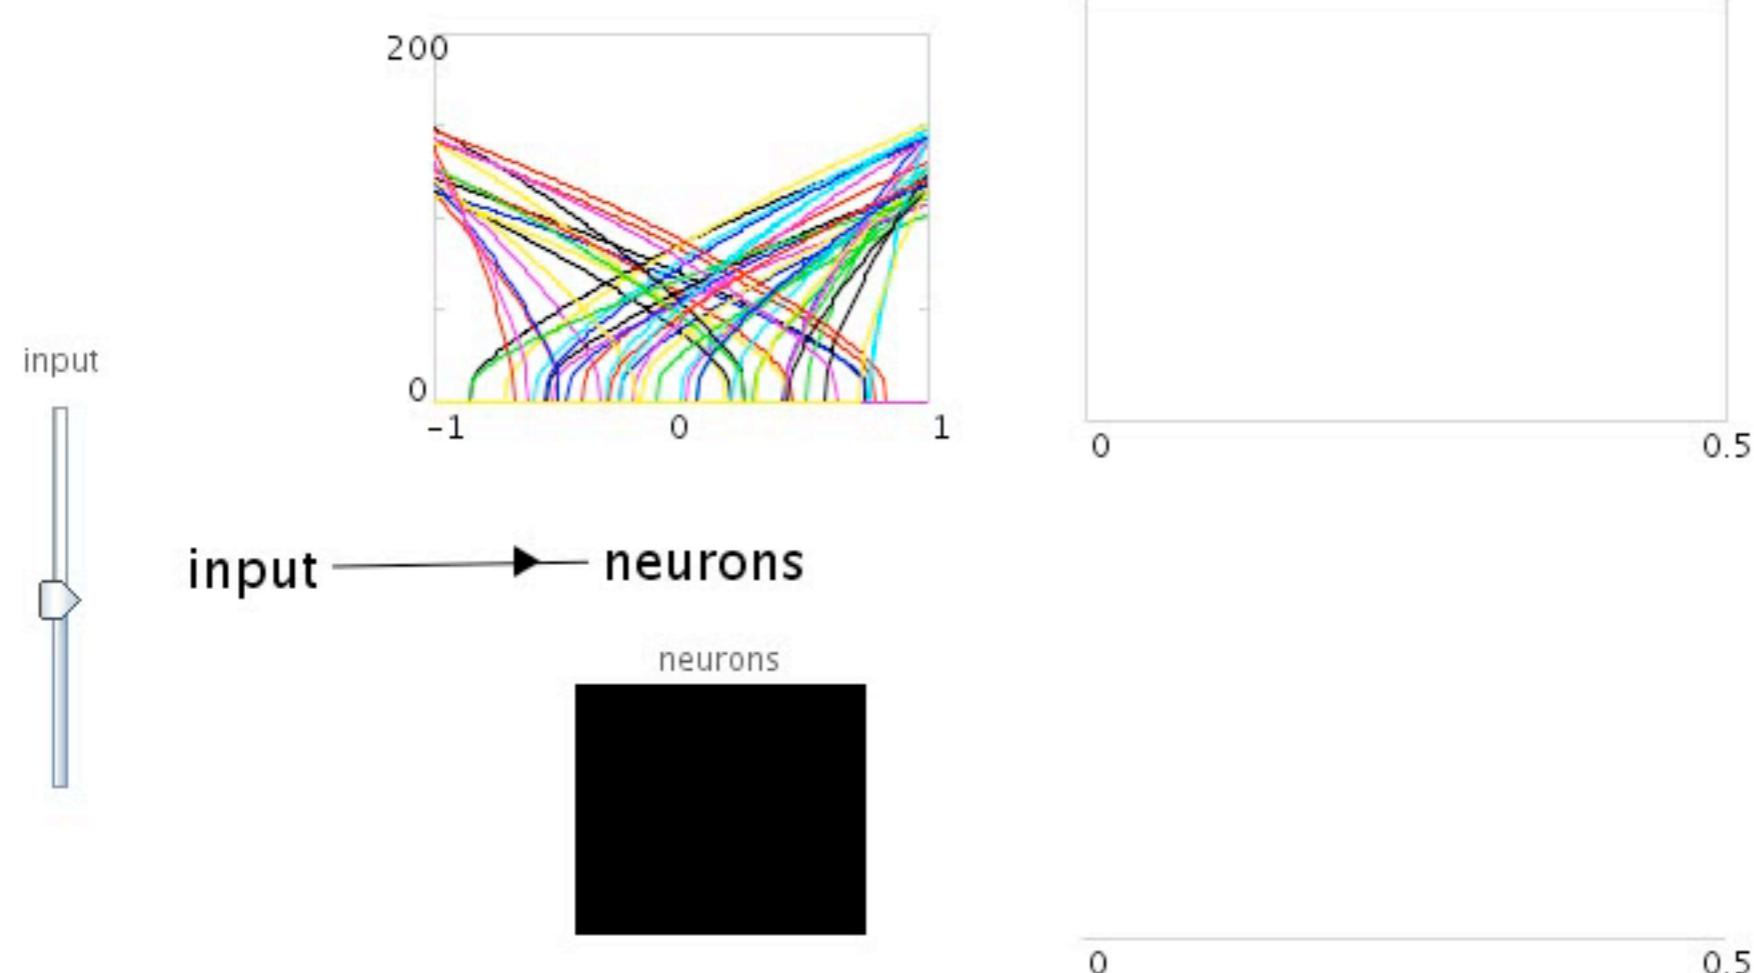
Principle I: Representation



$$a_i = G_i [\alpha_i e_i x + J_i^{bias}]$$
$$\hat{x} = \sum_i a_i \mathbf{d}_i$$

N = do sim

Principle I: Representation



$$a_i = G_i [\alpha_i e_i x + J_i^{bias}]$$
$$\hat{x} = \sum_i a_i \mathbf{d}_i$$

Learn the decoders

- Use the PES (prescribed error sensitivity) rule to learn the representation
 - Basically a delta rule with explicit error:

$$\Delta d_i = \kappa a_i E$$

The diagram illustrates the PES learning rule. It shows the equation $\Delta d_i = \kappa a_i E$. Two curved arrows point to the term a_i : one from the term κ and another from the term E . The arrow from κ is labeled "learning rate" and the arrow from E is labeled "error".

$$(\Delta \omega_{ij} = \kappa(e_j E)a_i)$$

Or, even easier

- Surprisingly simple optimization:

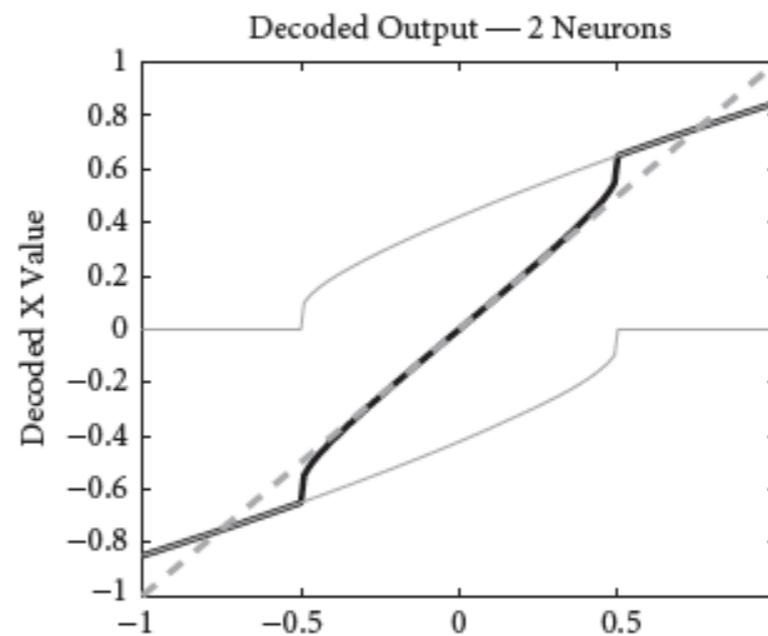
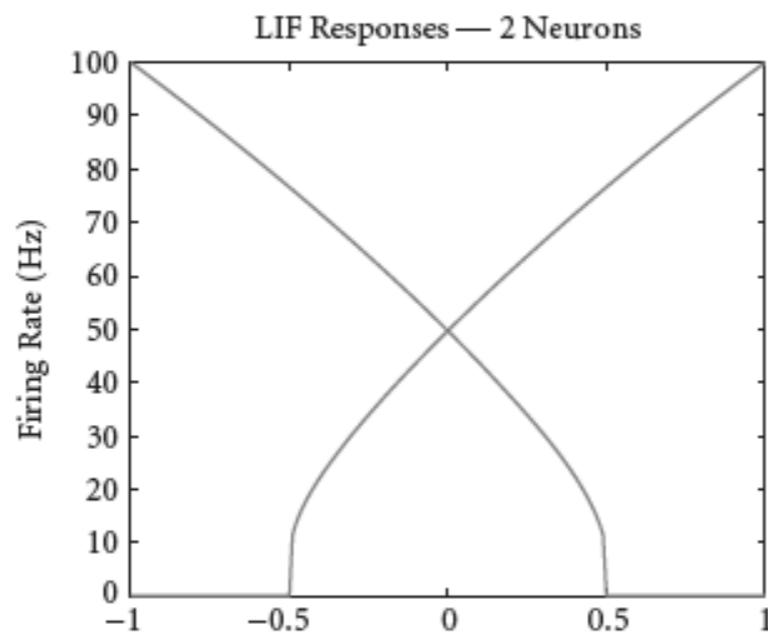
Q: How to find population ‘decoders’, \mathbf{d}_i

A: Minimize $\langle (x - \hat{x})^2 \rangle_x$ Recall: $\hat{x} = \sum_i a_i \mathbf{d}_i$

$$d^x = \Gamma^{-1} \Upsilon$$

$$= \left[\int_x a_i(x) a_j(x) dx \right]^{-1} \int_x x a_i(x)$$

d_i

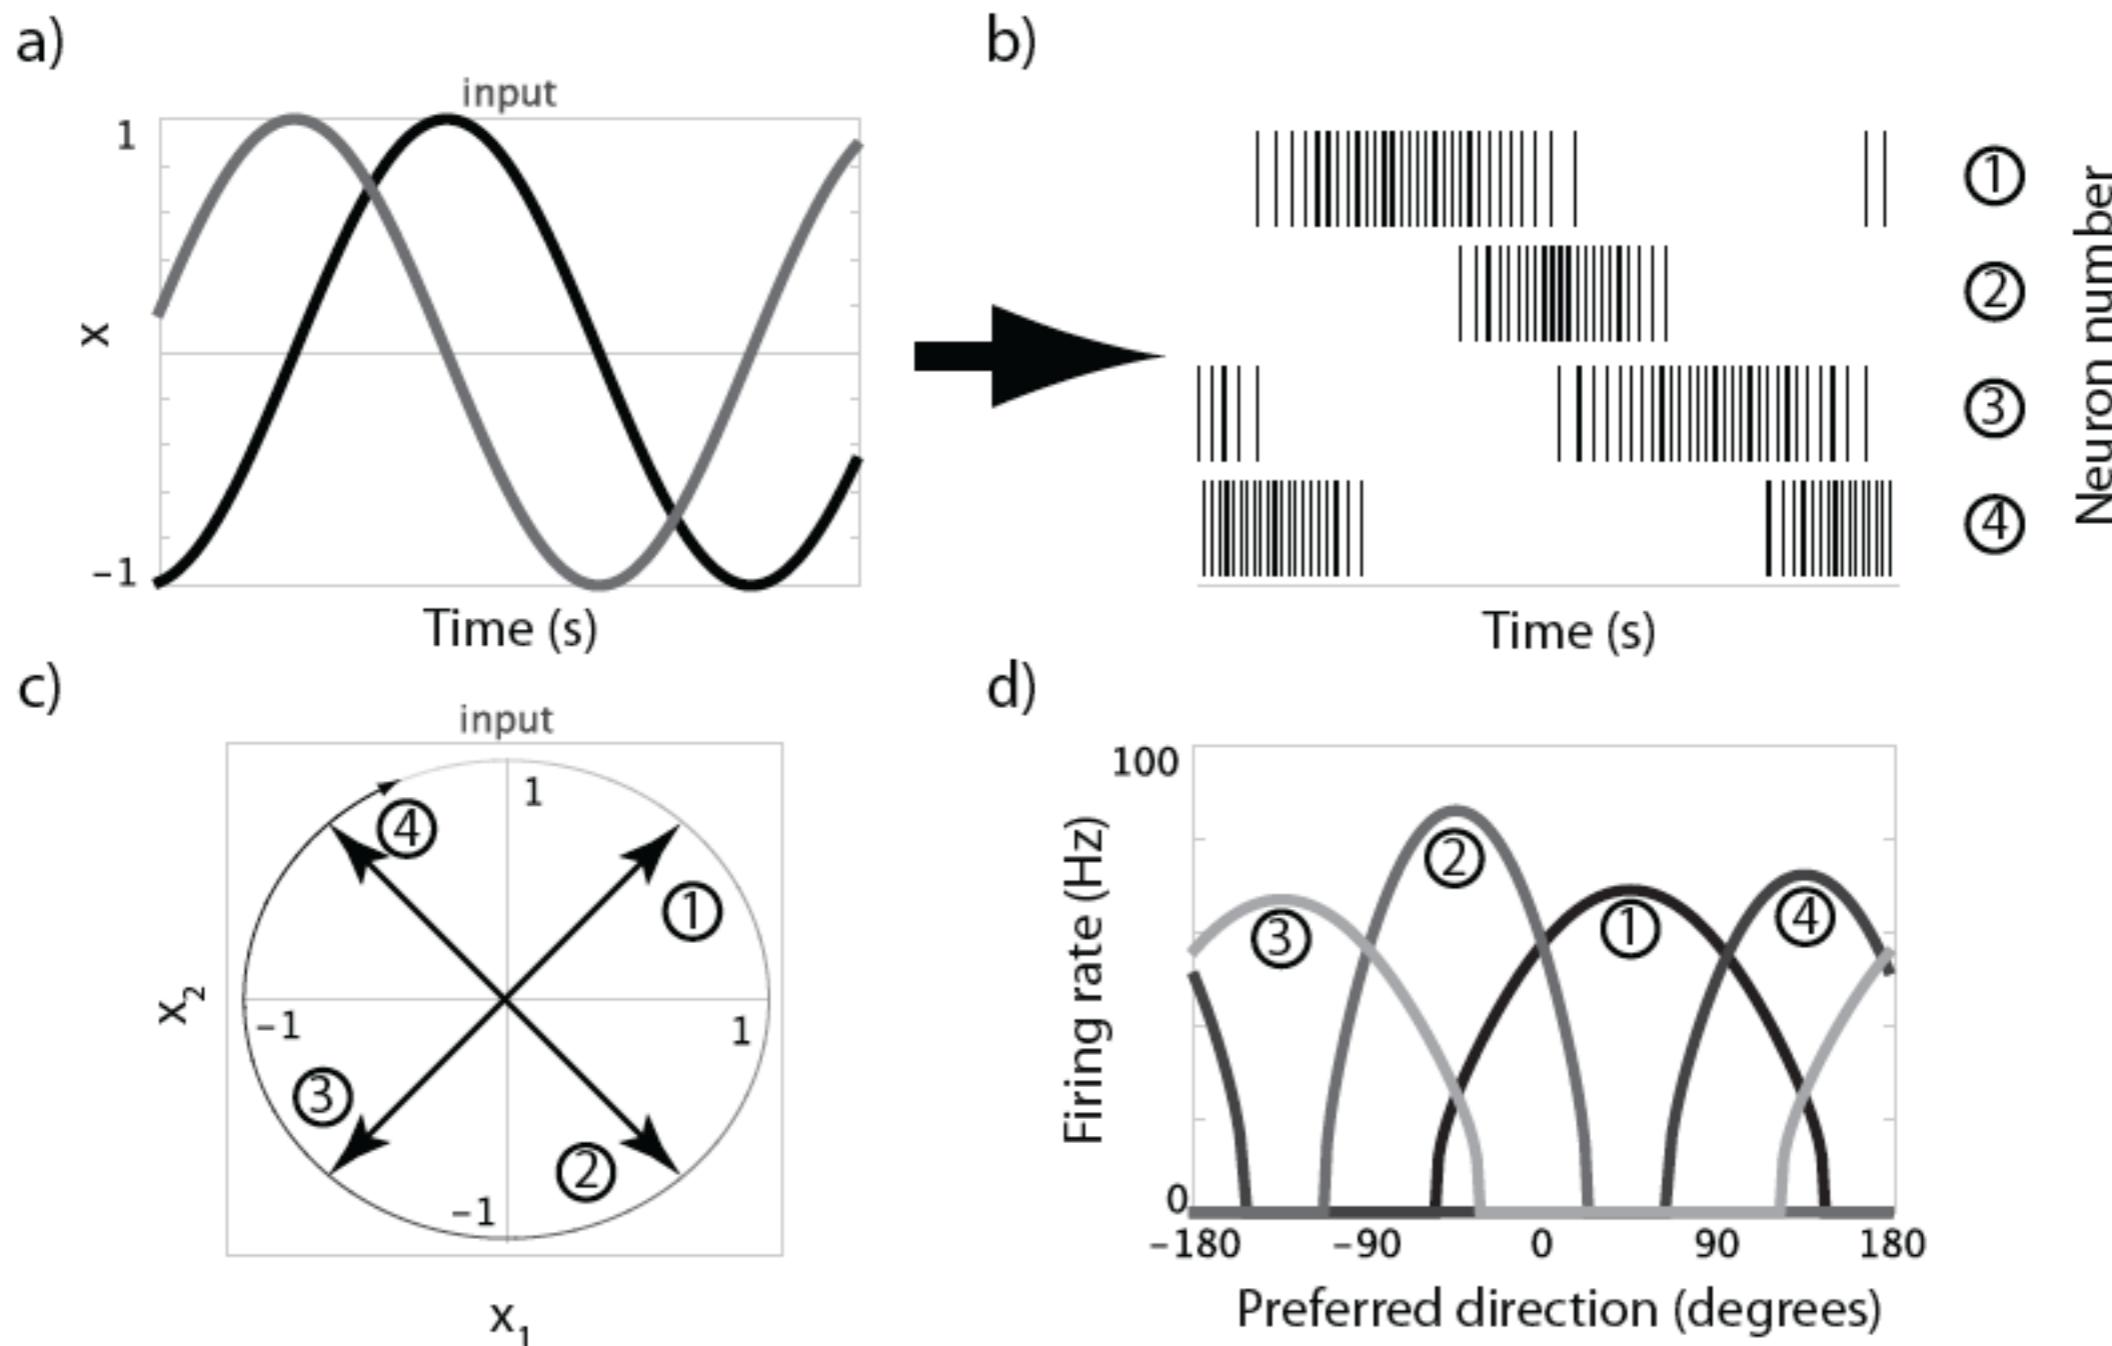


Input X Value

Input X Value

2D Representation

- ... it's the *same thing*



Representation Hierarchy

<i>Neural System</i>	<i>Dim</i>	<i>Encoder</i> ($a_i(x)$)	<i>Decoder</i> (\hat{x})
----------------------	------------	-----------------------------	------------------------------

NEF and Nengo

Centre for Theoretical Neuroscience
Nengo Summer School

Chris & Terry



Example: Working memory

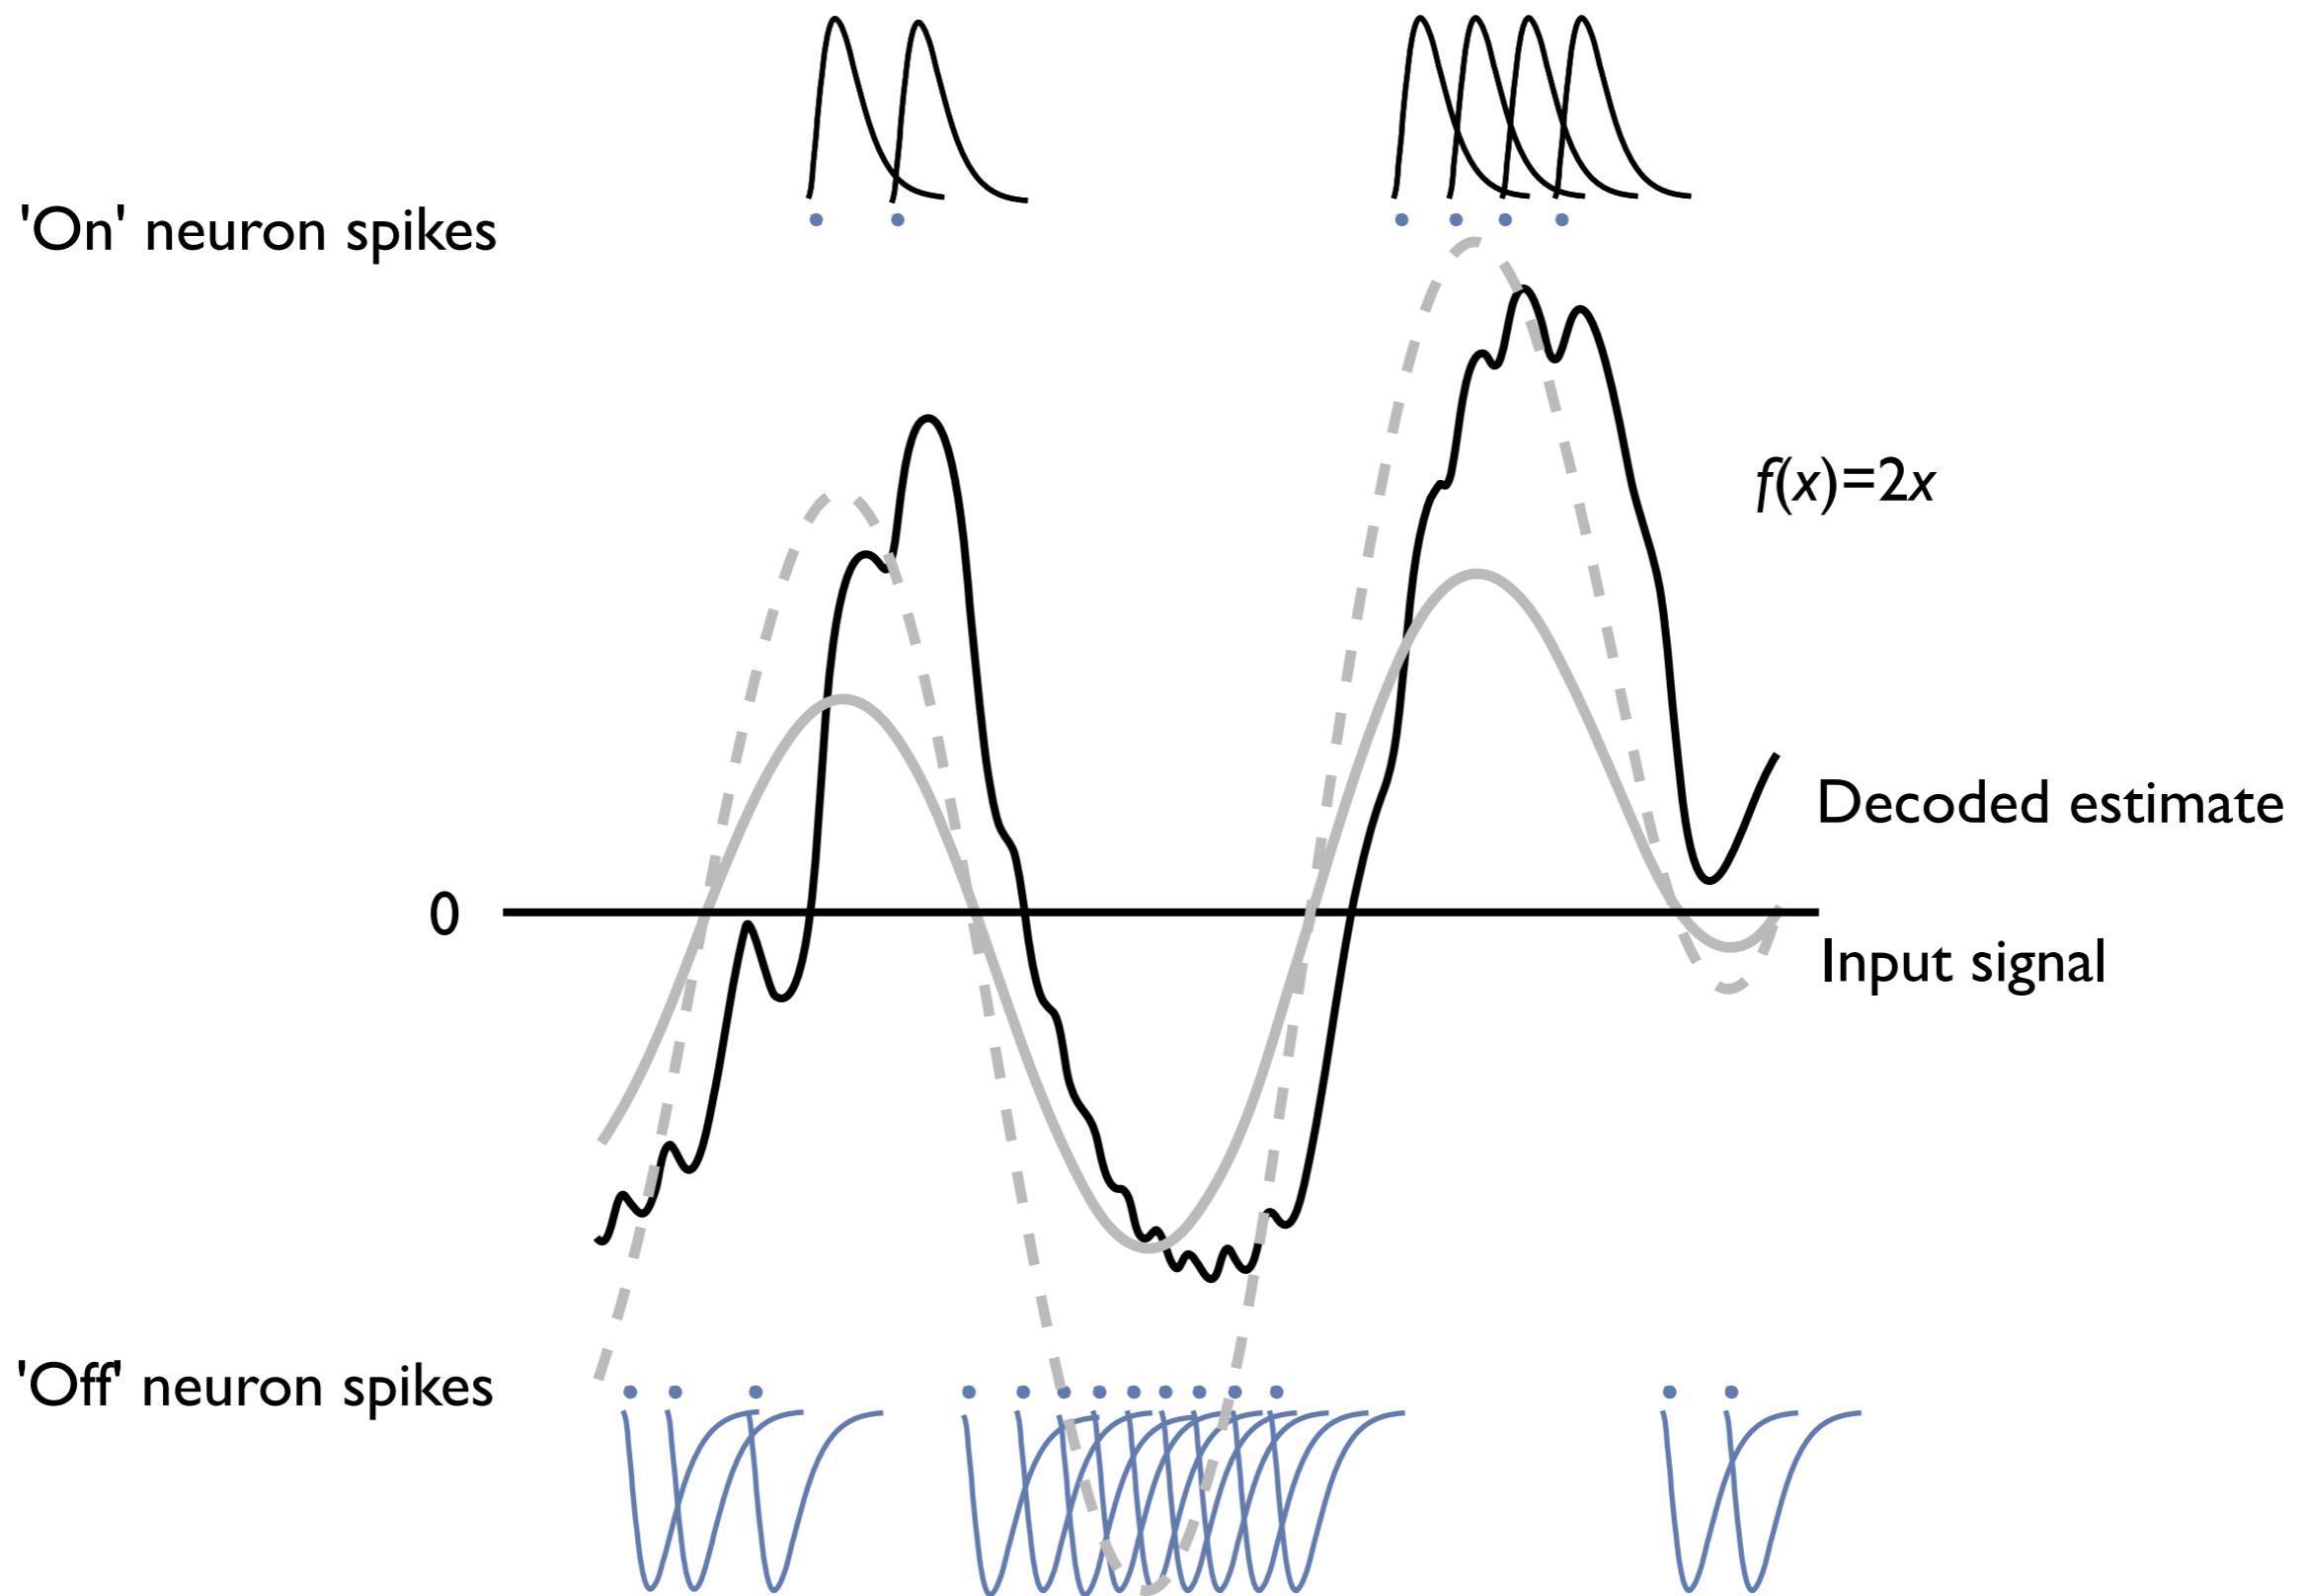
- We want a neural circuit that:
 - stores an input state
 - can be cleared (controlled storage)
- Dynamical equation (i.e., high-level program):

Memory $\dot{\mathbf{x}} = \alpha \mathbf{I} \mathbf{x}(t) + \mathbf{B} \mathbf{u}(t)$

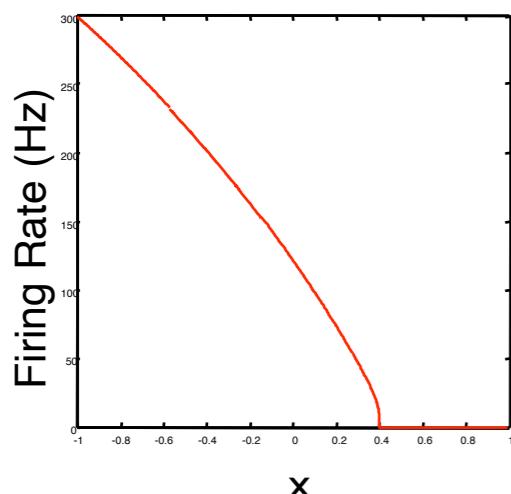
Input
Control

The diagram shows a circle labeled "Memory" containing the differential equation $\dot{\mathbf{x}} = \alpha \mathbf{I} \mathbf{x}(t) + \mathbf{B} \mathbf{u}(t)$. An arrow labeled "Input" points to the term $\mathbf{B} \mathbf{u}(t)$, and an arrow labeled "Control" points to the same term.

Principle 2: Transformation



Neuron Tuning Curve

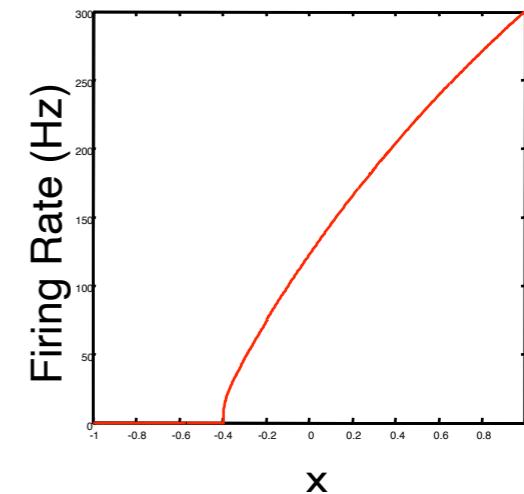


Left

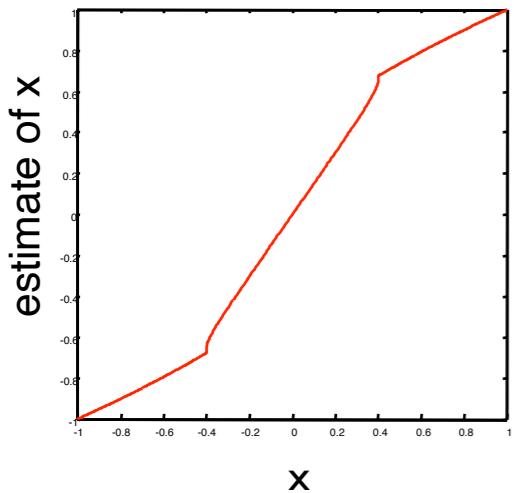
Encoding

Right

Neuron Tuning Curve

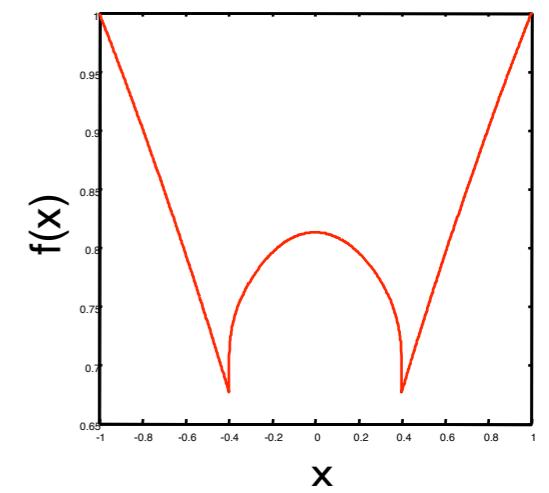


Decoding linearity



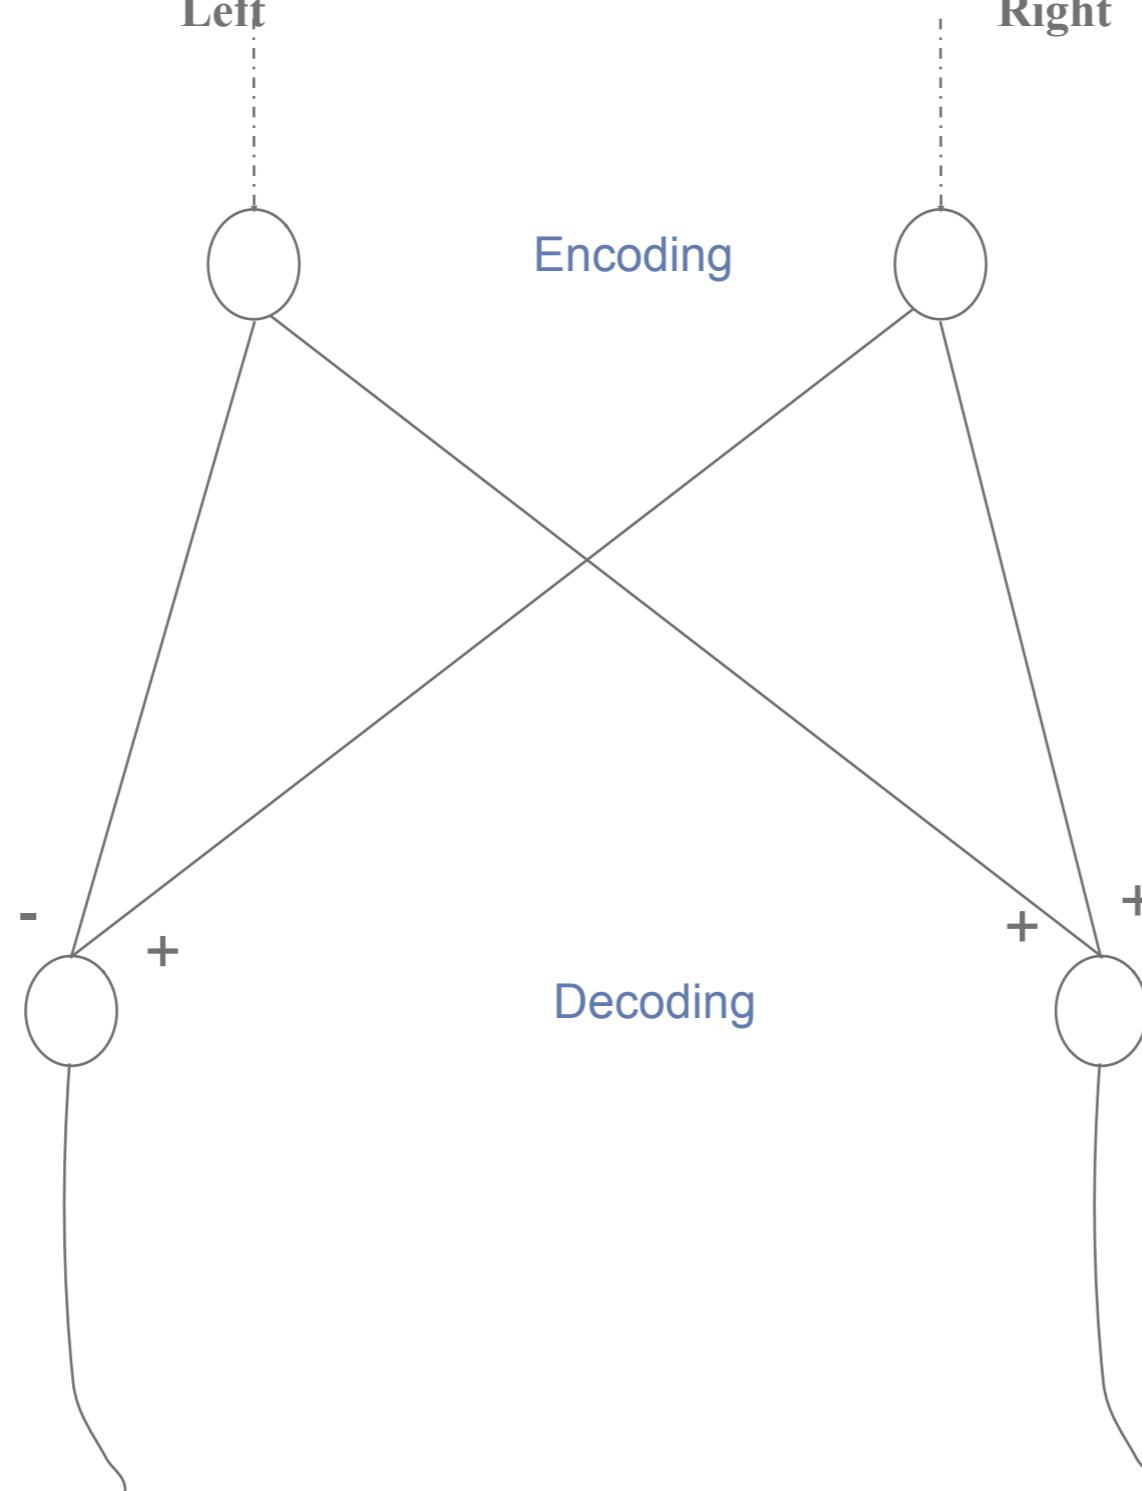
Decoding

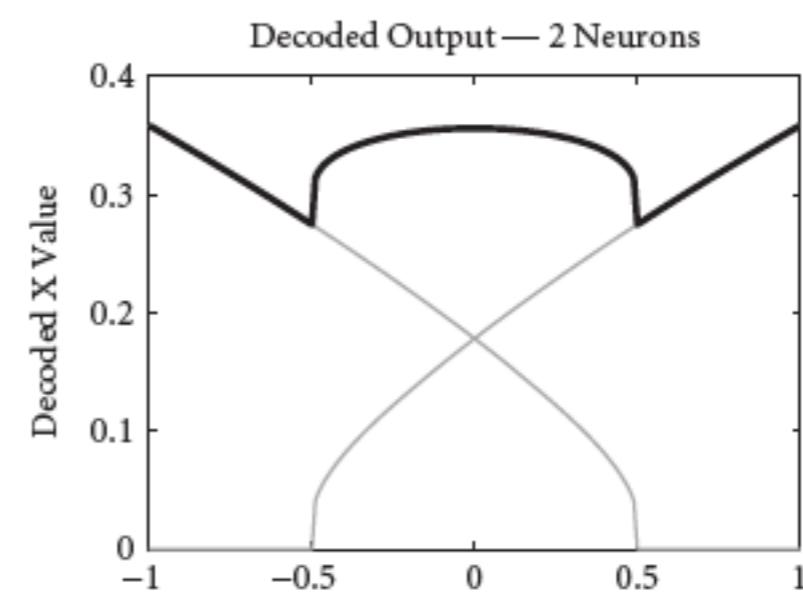
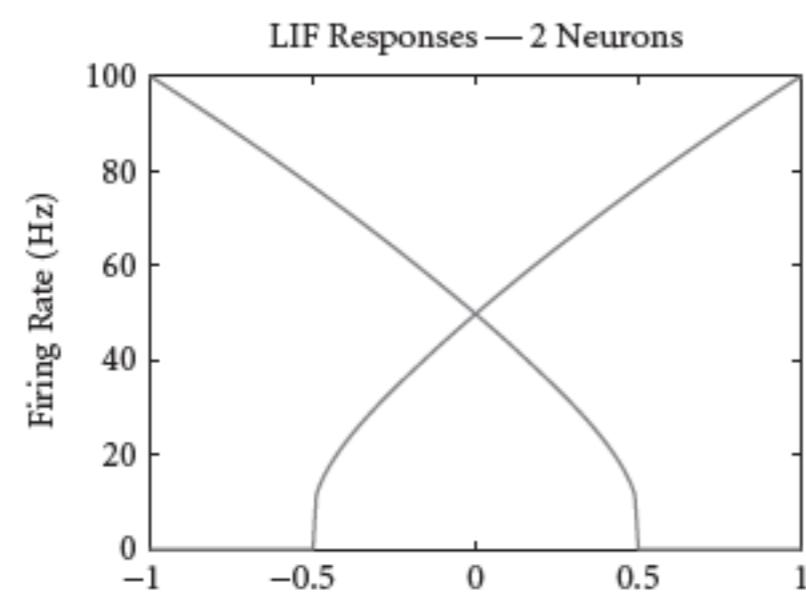
Decoding nonlinearity



Representation

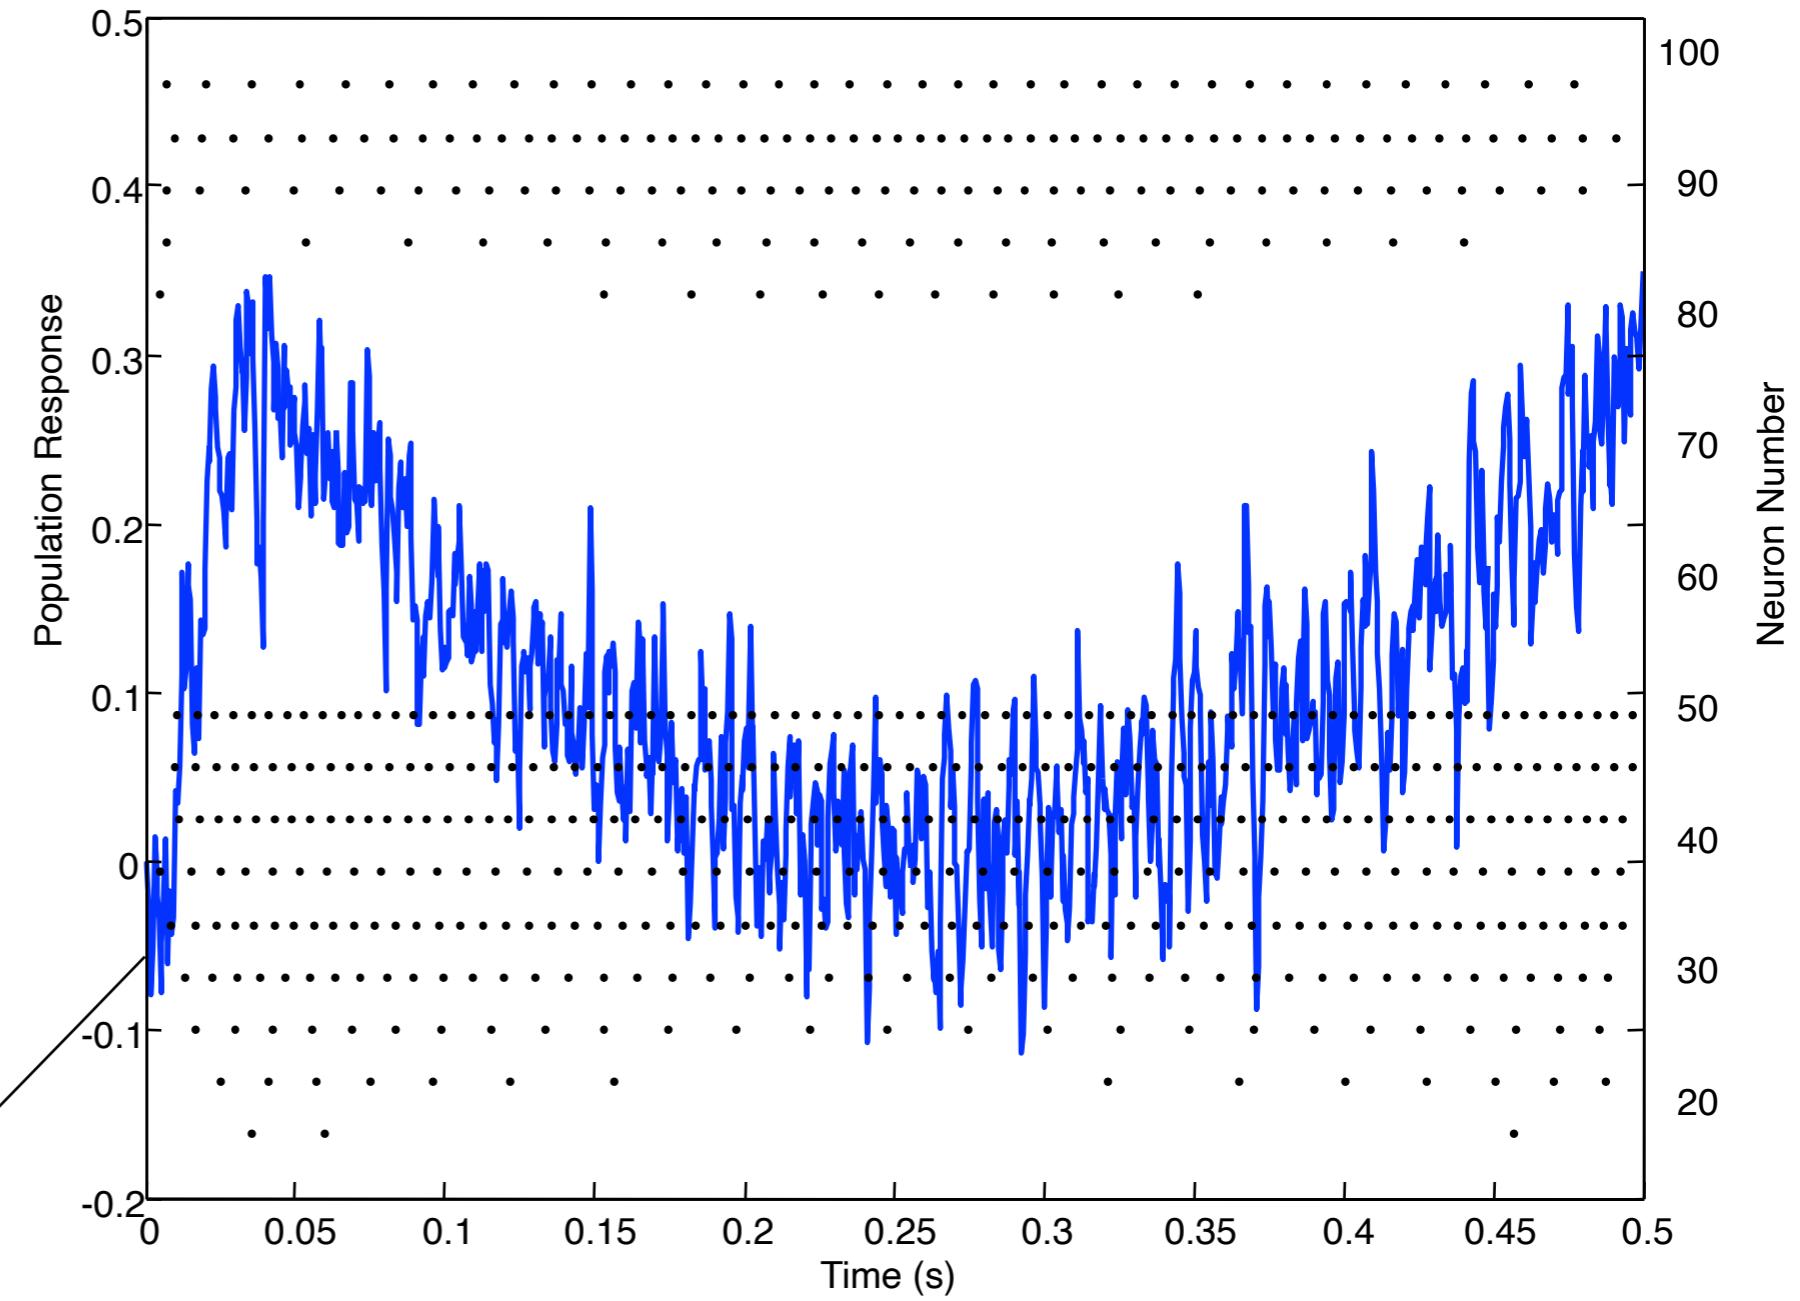
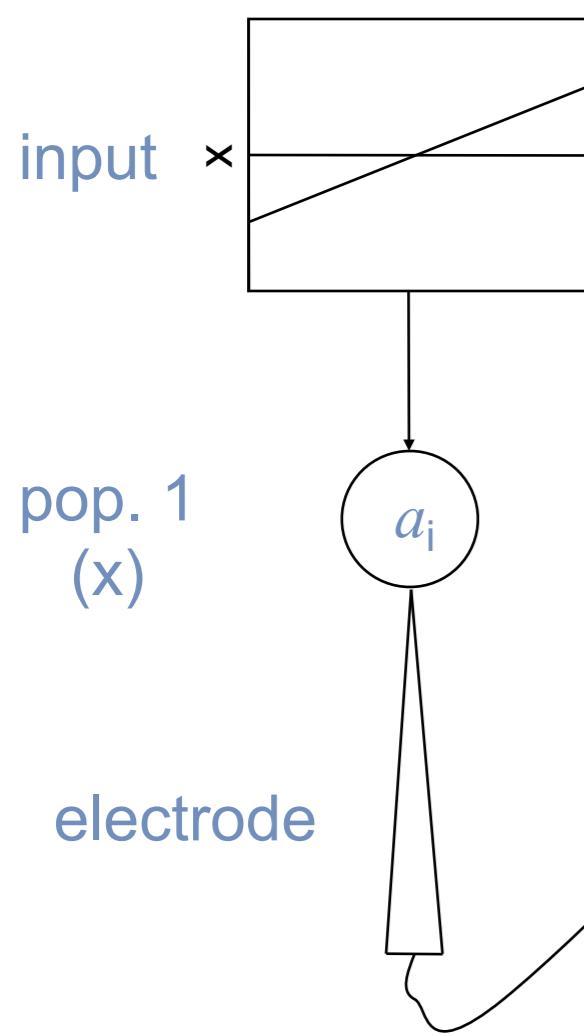
Transformation



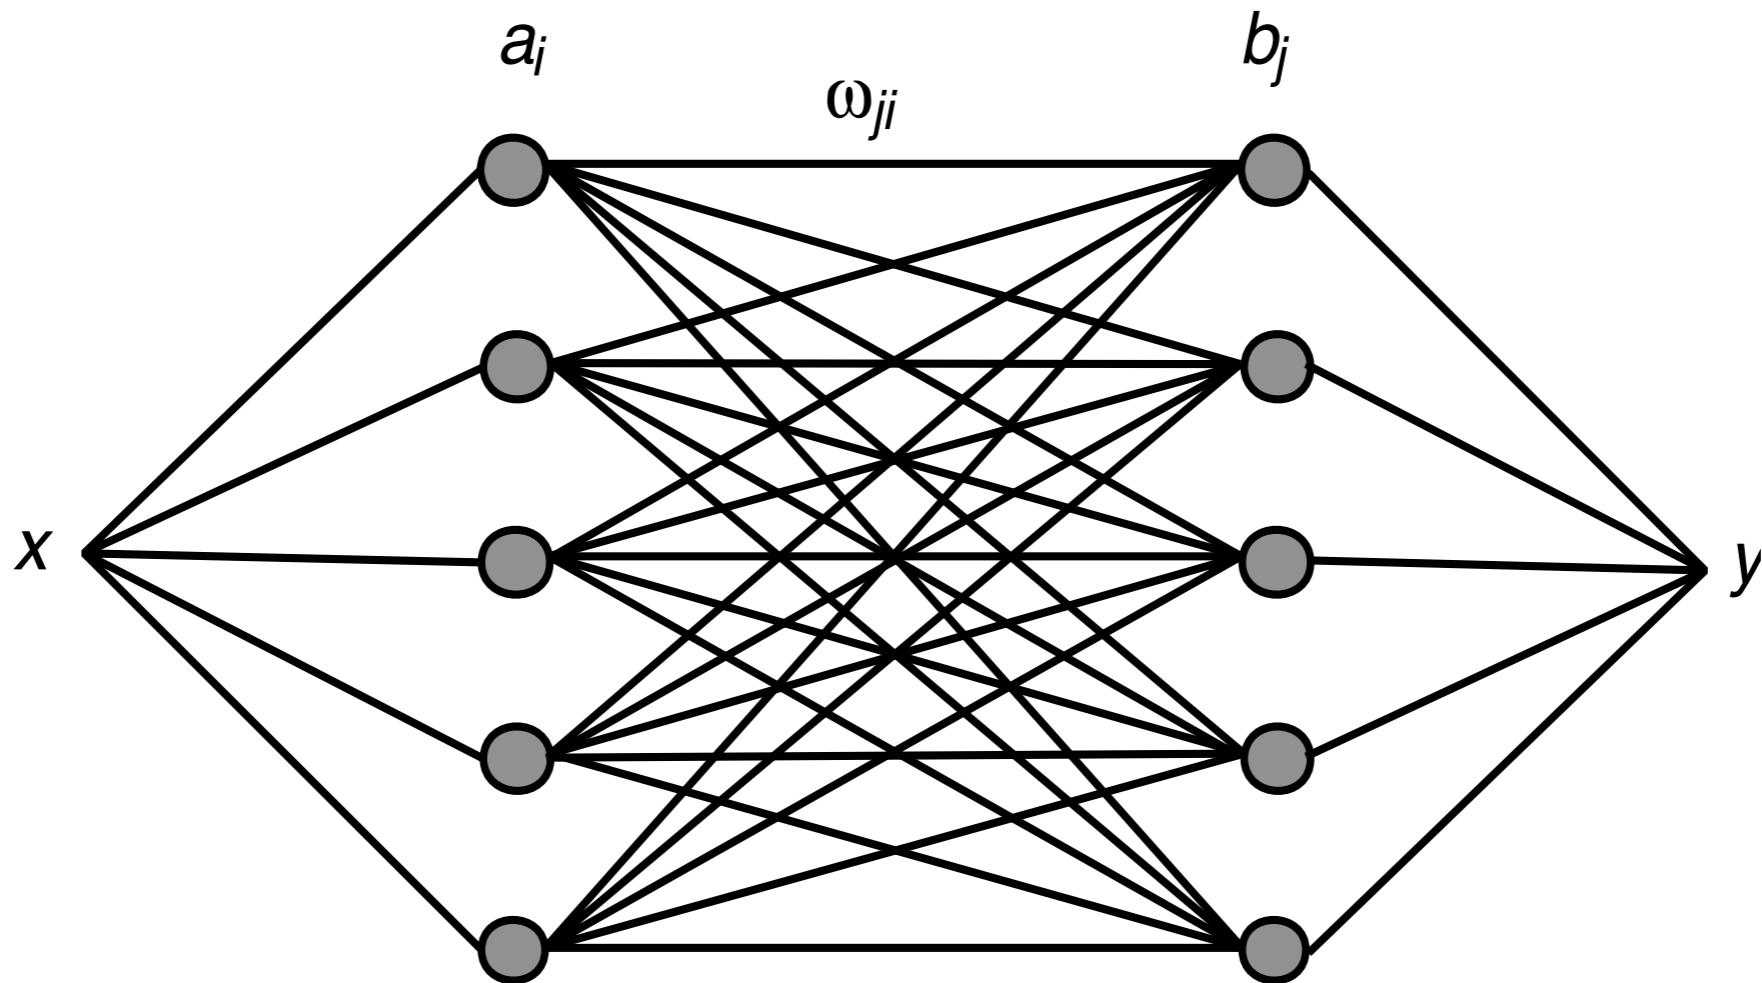


df_i

Principle 2: Computation



Building Circuits - A communication channel



- Define the representations of both pops:

Encoding

$a_i = G_i [J_i(x)]$	$b_j = G_j [J_j(y)]$
$= G_i [\alpha_i \mathbf{e}_i x + J_i^{bias}]$	$= G_j [\alpha_j \mathbf{e}_j y + J_j^{bias}]$

$\hat{x} = \sum_i a_i \mathbf{d}_i$	$\hat{y} = \sum_j b_j \mathbf{d}_j$
-------------------------------------	-------------------------------------

Decoding

- Define the computation: $y=x$
- Substitute our estimate of x into b

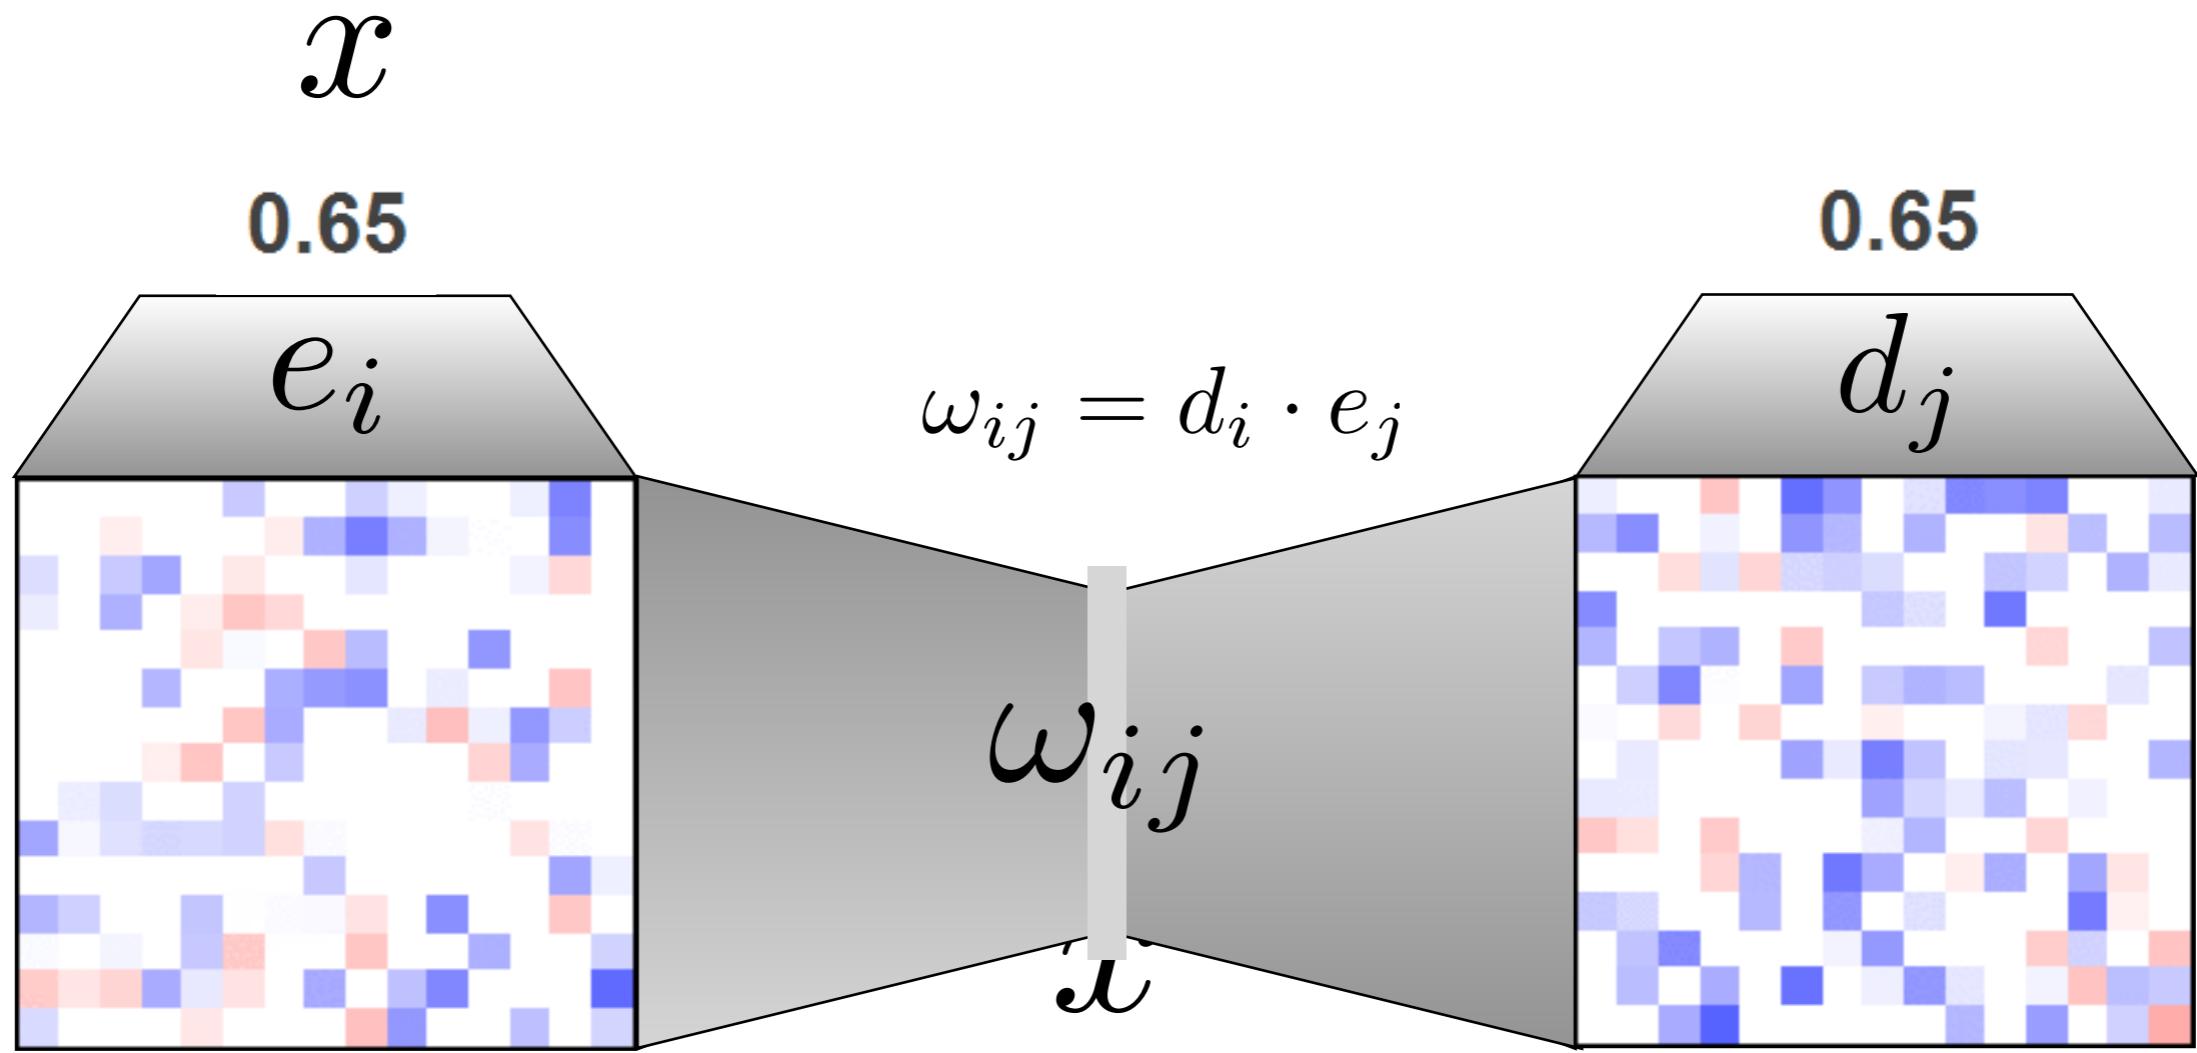
- Substituting:

$$y = x \approx \hat{x}$$

$$\begin{aligned} b_j &= G_j [\alpha_j \mathbf{e}_j x + J_j^{bias}] \\ &= G_j \left[\alpha_j \mathbf{e}_j \sum_i a_i \mathbf{d}_i + J_j^{bias} \right] \\ &= G_j \left[\sum_i \omega_{ji} a_i + J_j^{bias} \right] \end{aligned}$$

$$\omega_{ji} = \alpha_j \mathbf{e}_j \mathbf{d}_i$$

NEF connection weights



$$J_i = x \cdot e_i$$

$$a_i = G_i [J_i(x)]$$

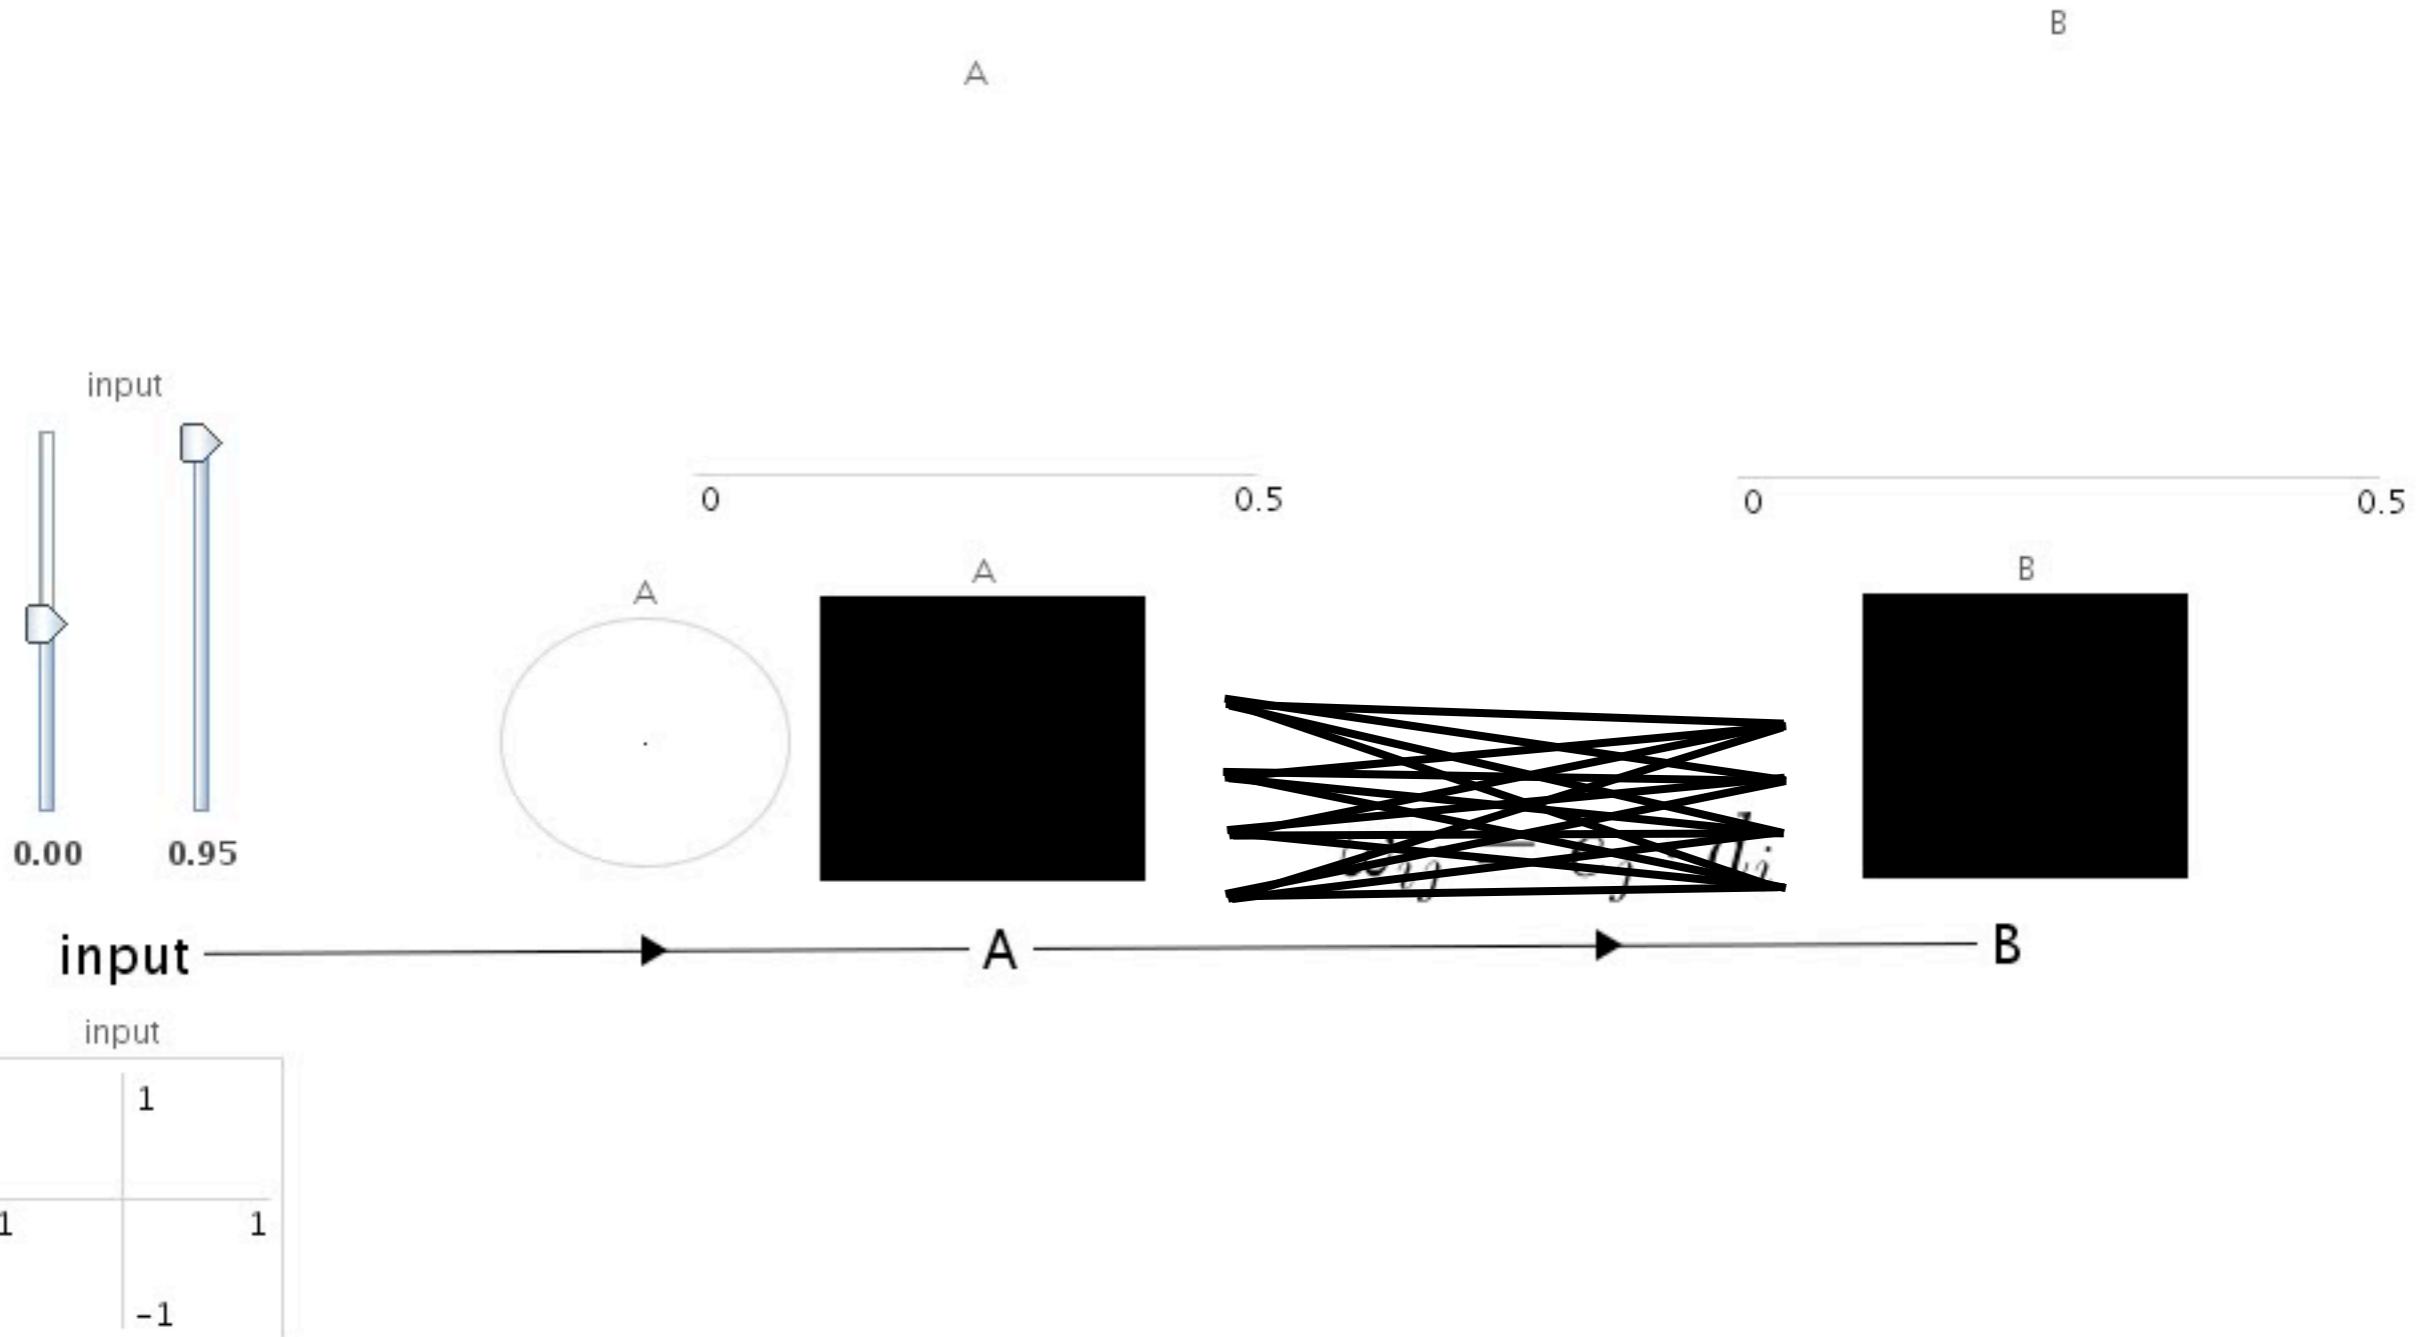
$$\hat{x} = \sum_i a_i d_i^x \leftarrow f(x)$$

$$J_j = \sum_i \omega_{ij} a_i$$

$$b_j$$

$$J_j = \hat{x} \cdot e_j$$

$$b_j = G_j [J_j]$$



Pick a function

- Any function... build a feedforward approximation

Now learn it

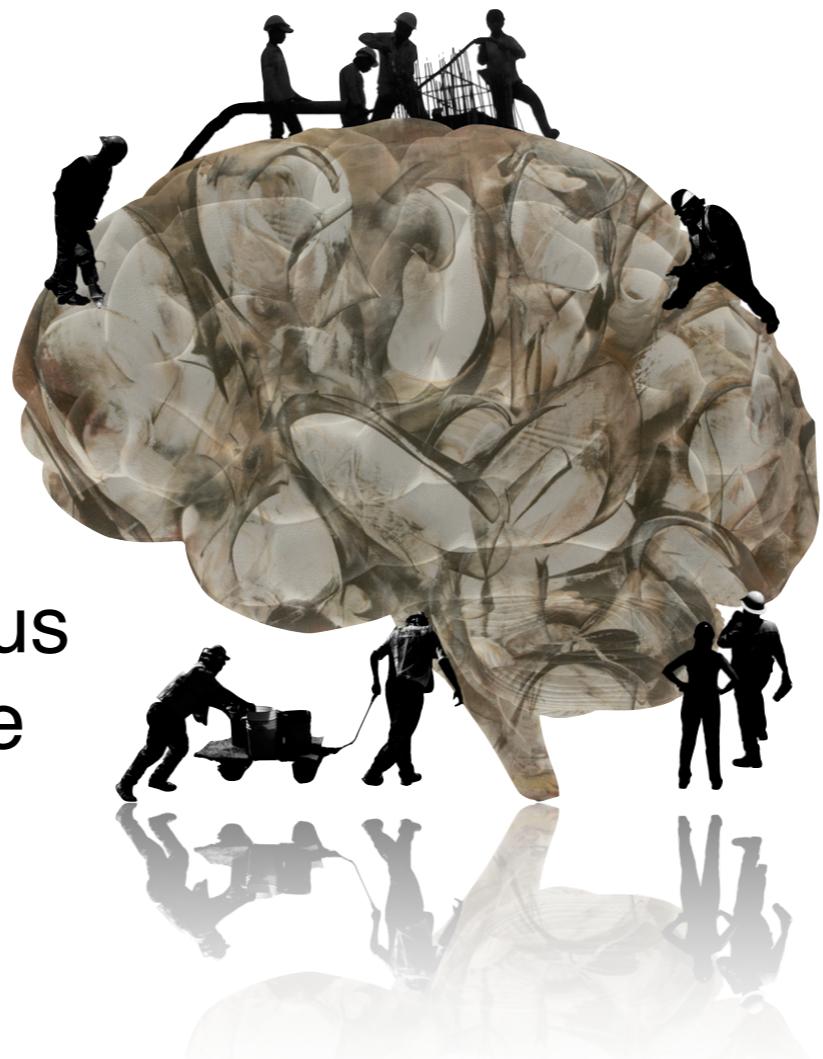
$$\Delta d_i = \kappa a_i E$$

The diagram illustrates the components of the update rule. The equation $\Delta d_i = \kappa a_i E$ is displayed above. Below it, two curved arrows point upwards from the words "learning rate" and "error" to the terms κa_i and E respectively.

$$(\Delta\omega_{ij} = \kappa(e_j E)a_i)$$

Summary

- Built a feedforward computation in spiking neural networks
- Principles used are very general (vector space representation, nonlinear computation)
- Dealt with heterogeneity, nonlinearities, noise
- ... dynamics & scaling... next



Previous
Lecture

Next
Lecture

Playlist

Centre for Theoretical Neuroscience
Nengo Summer School

Chris & Terry